

Date:

Task No: 1

Design a multi-dimensional data model schema namely Star, Snowflake and Fact Constellations for a Categorical data using SQL Server Management Studio (SSMS). (Perform the above for Banking, Healthcare, Manufacturing, Sales and Automobile)

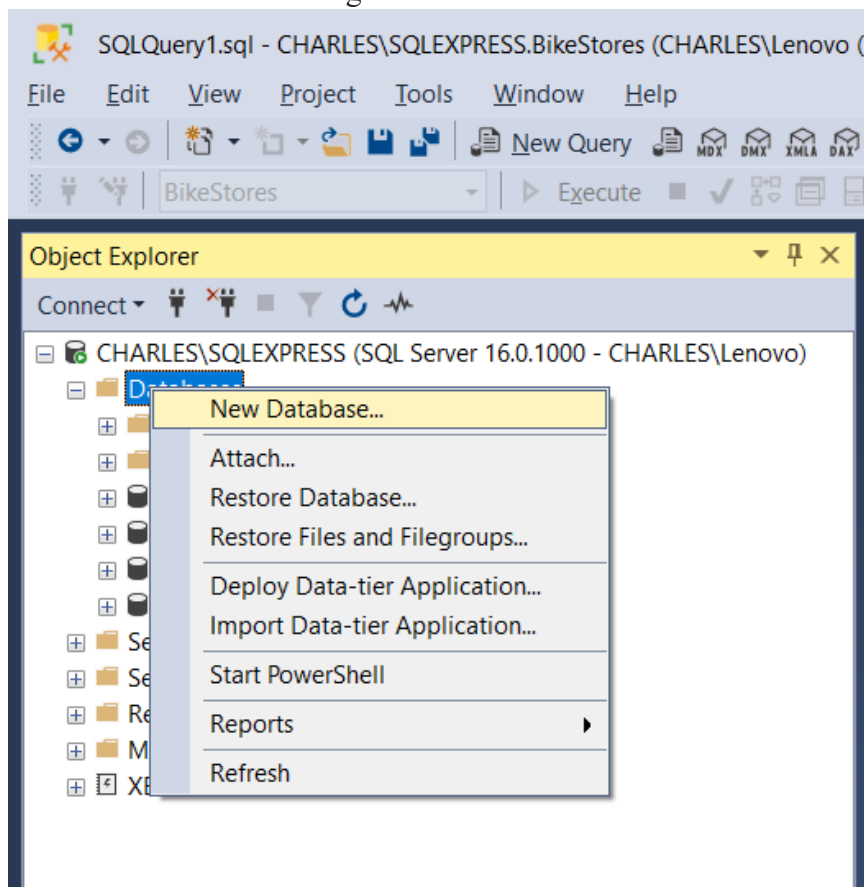
AIM:

PROCEDURE:

To practice creating a star schema data model from scratch, install SQL Server Management Studio, Microsoft SQL Server

Step 1: Install Diagram Support.

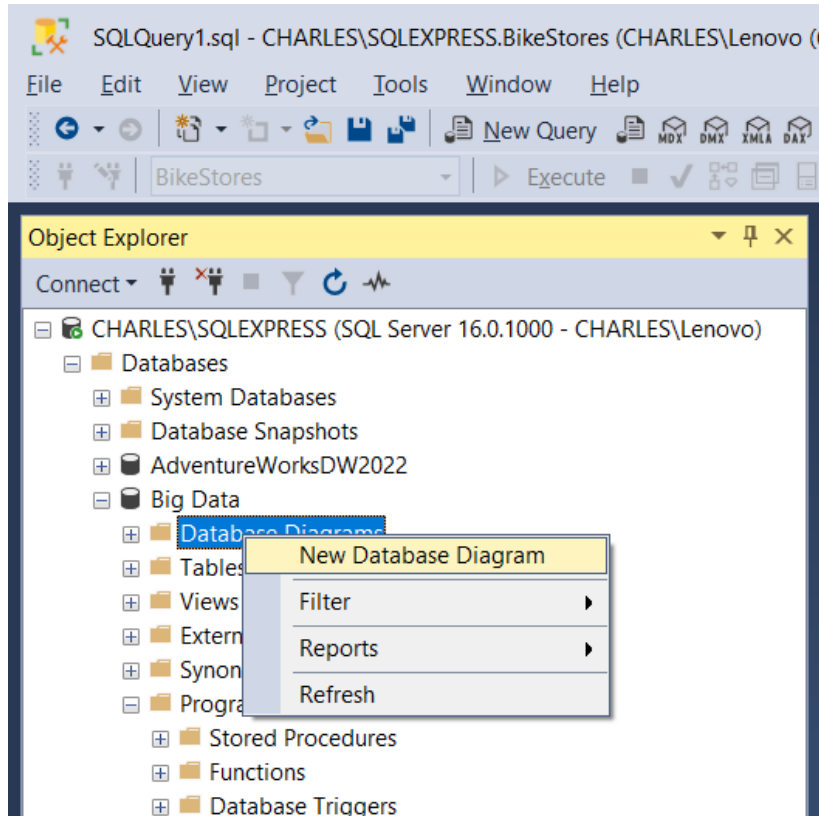
- Create a new database “BigData”



- Databases -> BigData -> Database Diagrams > Install Diagram Support
- Select the “Yes” in the pop-up window to close the window to one or more create support objects

Step 2: Create New Database Diagram.

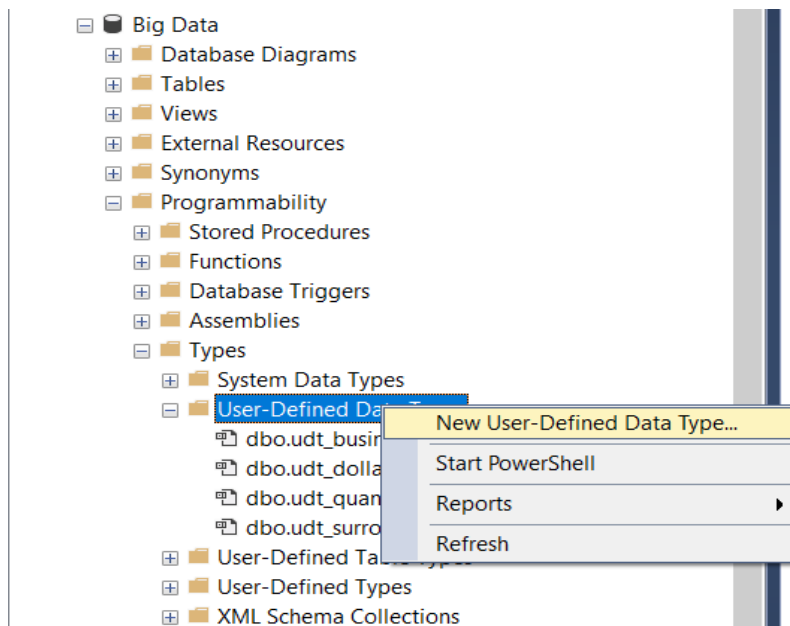
- Right-click on the menu item “Database Diagrams” and select the “New Database Diagram” item.



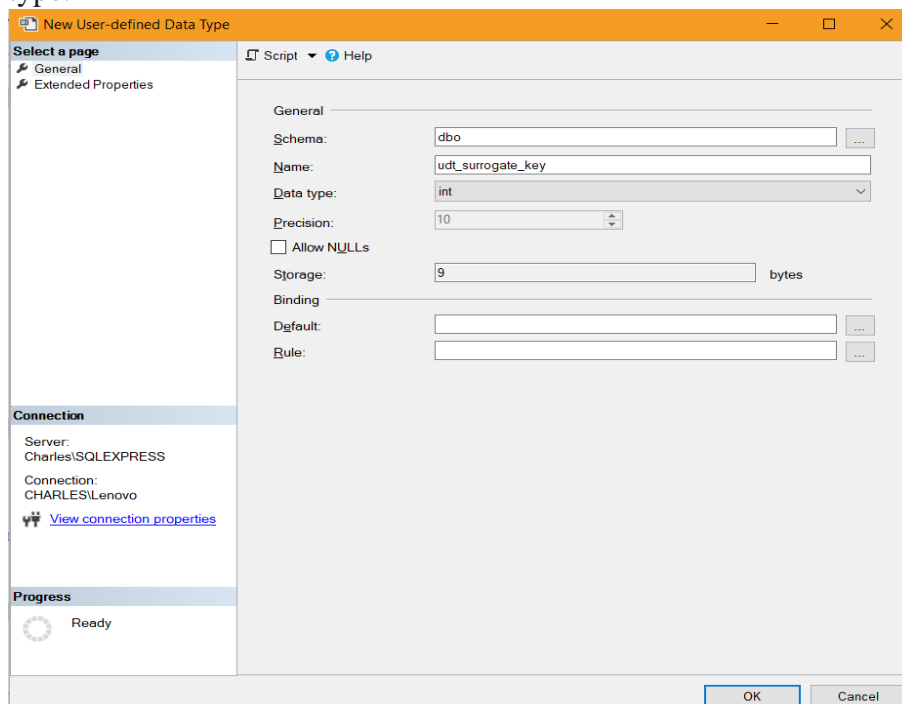
- A pop-up window with a title “Add Table” shows up. Click on the “Close” button at the bottom of the new window to close the window

Step 3: Create User-Defined Data Types

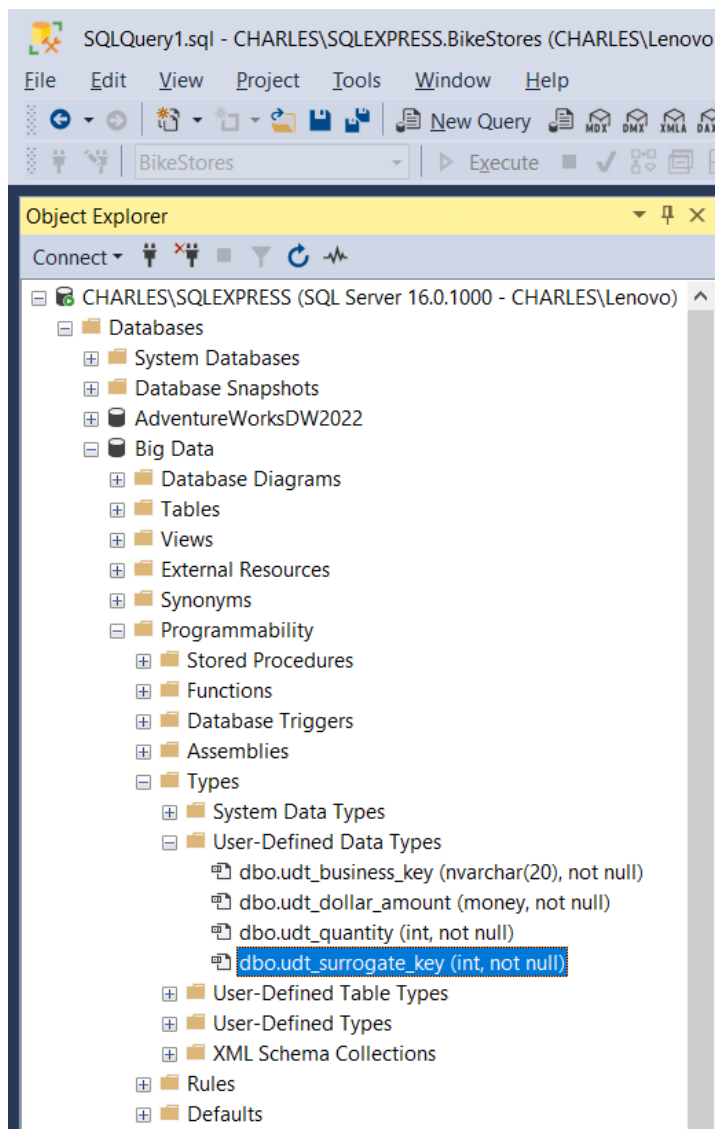
- Right-click on the item “Databases -> data_modeling -> Programmability -> Types -> User-Defined Data Types” in the Object Explore panel and select the menu item “New User-Defined Data Type” in the context menu.



- A new window appears. Enter a value “udt_surrogate_key” in the name field and select the data type with “int”. Then, click on the “OK” button to create a new data type.



- Follow the preceding procedure described in this step, create other user-defined data types.

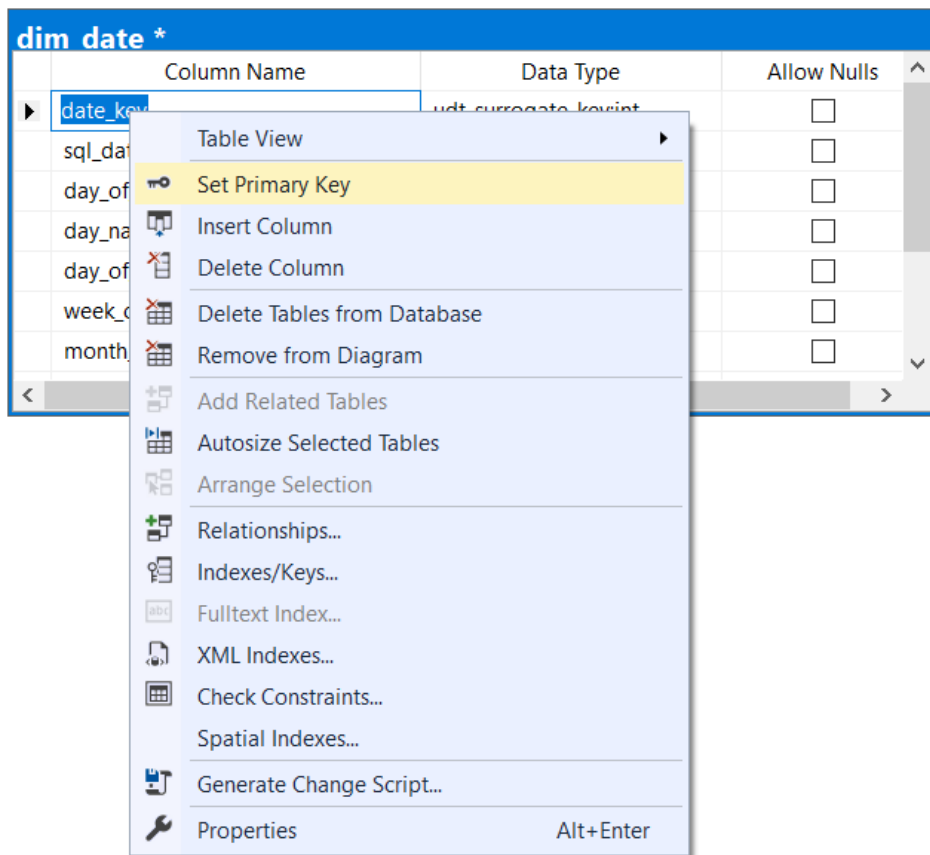


Step 4: Create a Dimension Table in SSMS.

- Right-click on an empty space in the middle panel, select the “New Table...” menu item in the context menu
- A small pop-up window shows up and enter a value “dim_date” as the table name. Click the “OK” button to close this window and start a new pop-up window.
- In the new window, enter a value “date_key” as a column name and select the user-defined data type “udt_surrogate_key” as the data type. Then uncheck the “Allow Nulls” checkbox.

dim_date			
	Column Name	Data Type	Allow Nulls
🔑	date_key	udt_surrogate_key:int	<input type="checkbox"/>
	sql_date	date	<input type="checkbox"/>
	day_of_week	int	<input type="checkbox"/>
	day_name_of_week	nchar(10)	<input type="checkbox"/>
	day_of_month	nchar(10)	<input type="checkbox"/>
	week_of_year	nchar(10)	<input type="checkbox"/>
	month_of_year	nchar(10)	<input type="checkbox"/>

- Right-click on trigonal icon beside the text “date_key”, select the “Set Primary Key” menu item in the context menu.

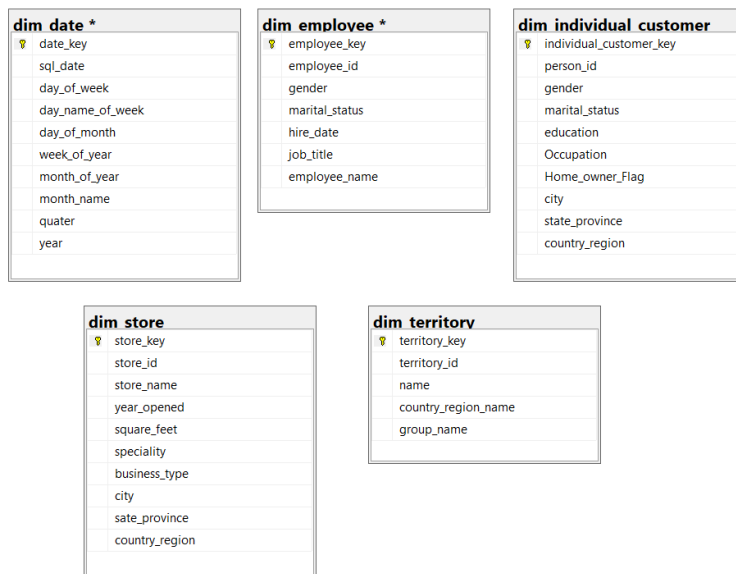


Step 5: Save the New Diagram.

- Click on the “Save” button on the window’s toolbar.
- A pop-up window shows up to ask a name of the diagram. Enter a value “fact_sale_order” as the name for the diagram then click on the “OK” button.

Step 6: Create All Dimension Tables.

- Repeat Step 5 to create other dimension tables.



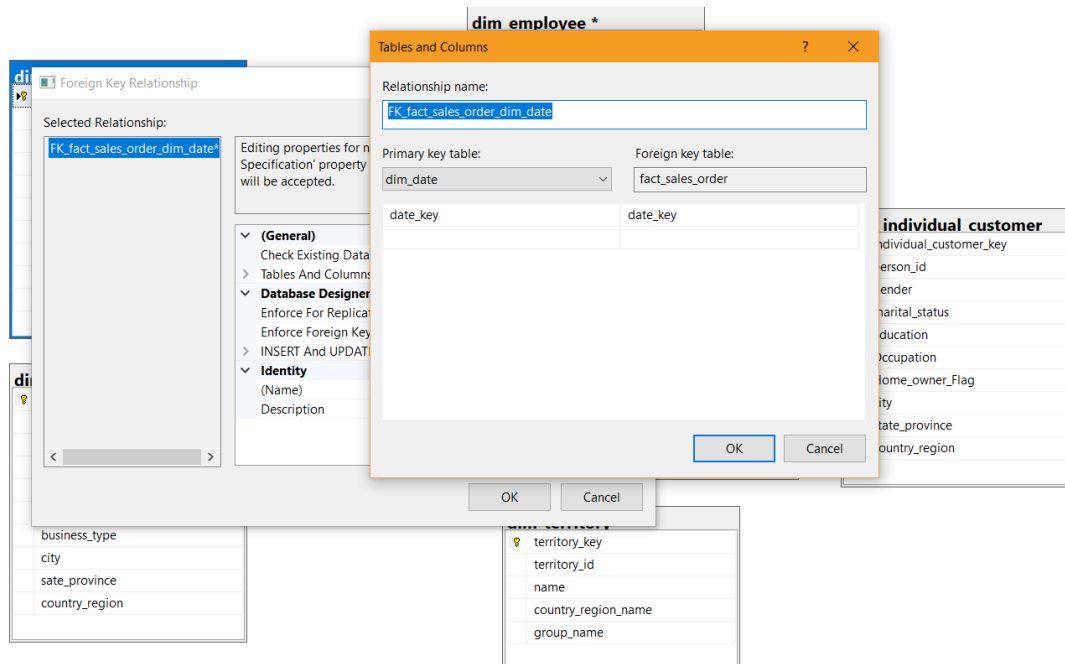
Step 7: Create a Fact Table.

- Right-click on an empty space in the diagram panel, select the “New Table...” menu item from the context menu
- enter a value “fact_sales_order” as the fact table name in the new pop-up window.
- Click the window by clicking on the “OK” button.
- The second pop-up window appears. Enter all dimension table keys and select the user-defined data type “udt_surrogate_key” as their data types. Uncheck all “Allow Nulls” checkboxes.

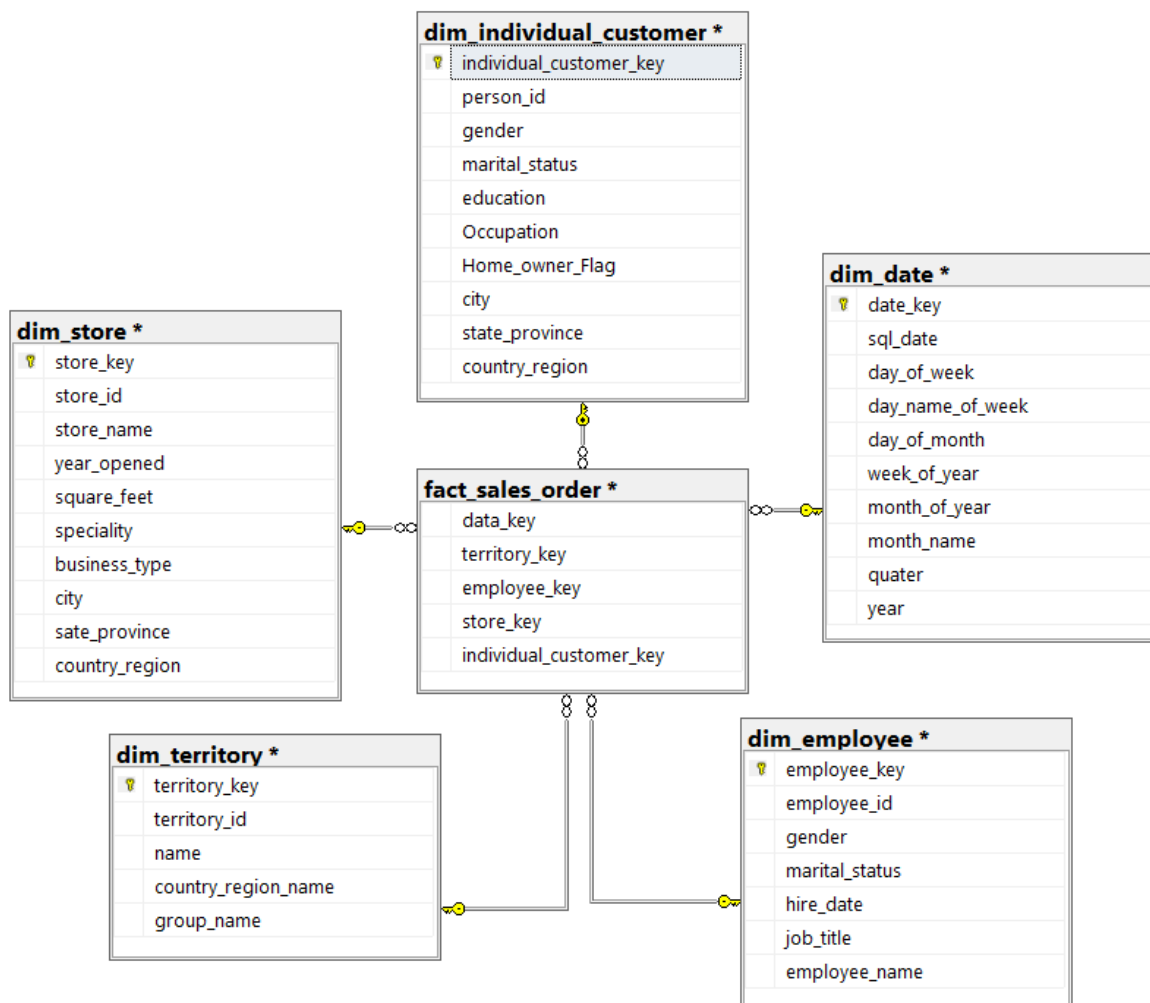
fact sales order *			
	Column Name	Data Type	Allow Nulls
	date_key	udt_surrogate_key:int	<input type="checkbox"/>
	employee_key	udt_surrogate_key:int	<input type="checkbox"/>
	store_key	udt_surrogate_key:int	<input type="checkbox"/>
	territory_key	udt_surrogate_key:int	<input type="checkbox"/>
	individual_customer_key	udt_surrogate_key:int	<input type="checkbox"/>
			<input type="checkbox"/>

Step 8: Establish Relationship Between Dimension tables and Fact table

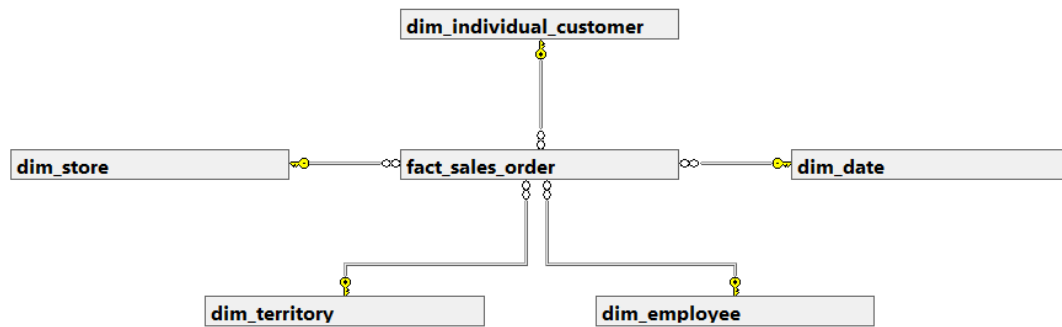
- In the diagram panel, drag the key icon in the table “dim_table” over the “date_key” in the table “fact_sales_order”. Two pop-up windows appear immediately.
- In the first pop-up window presents the relationship name, “FK_fact_sales_order_dim_date”. All the foreign names will use this format.
- verify if the keys are linked correctly and click OK button



Charles\SQLEXPRESS...- fact_sale_order* - X



Charles\SQLEXPRESS...- fact_sale_order* X



RESULT:

Date:

Task No: 2

To configure, monitor, and administer a Data warehouse and perform basic Query operations on the DW.

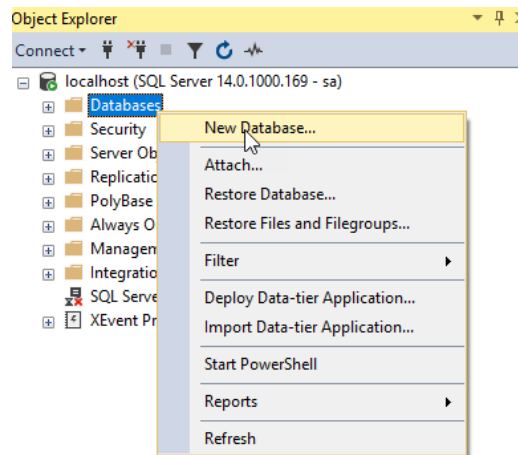
AIM:

TASK 2a:

PROCEDURE:

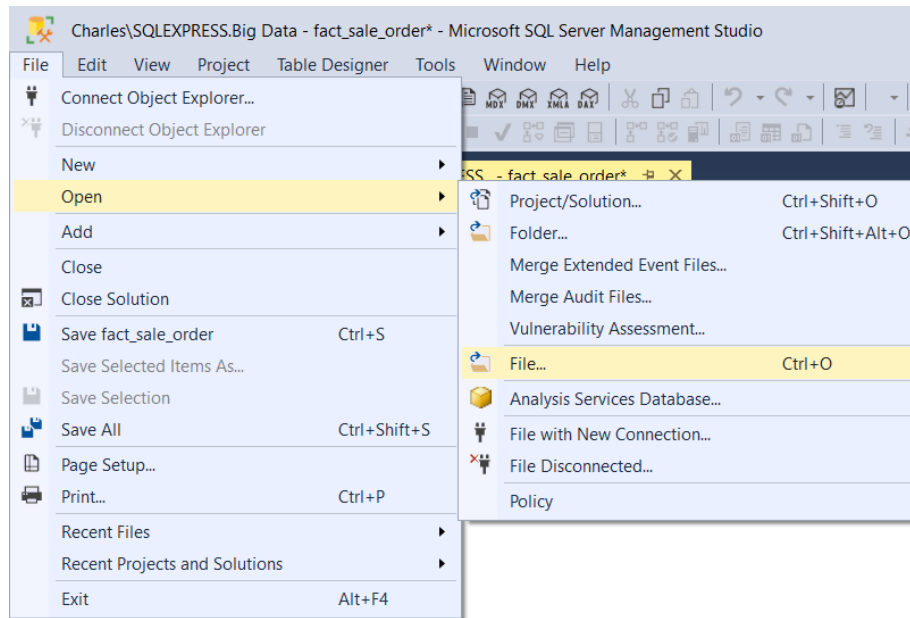
Connect to the SQL Server by choosing the server name, enter the user and password and click the Connect button.

Right-click the Databases node in the Object Explorer and select the New Database... menu item

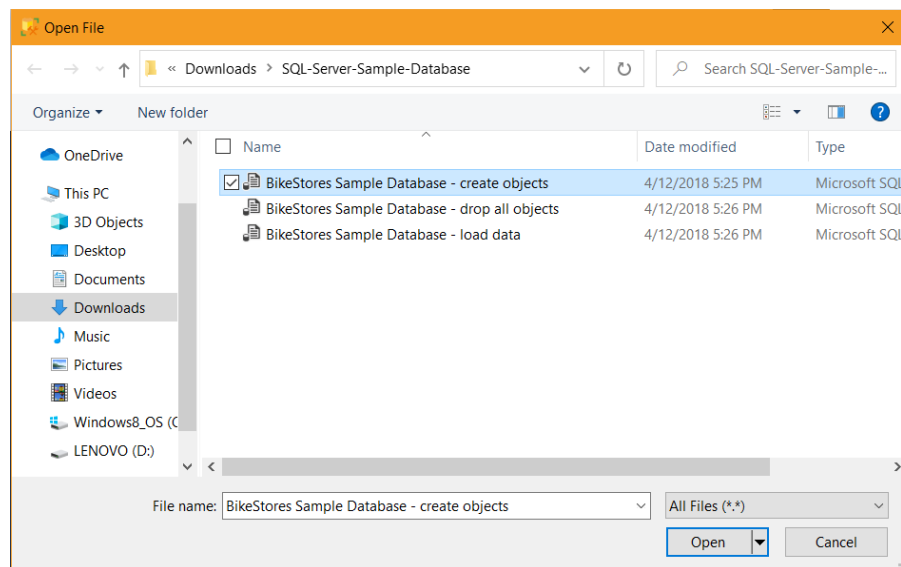


Enter the **Database name** as BikeStores and click the **OK** button to create the new database.


From the File menu, choose Open > File... menu item to open a script file.



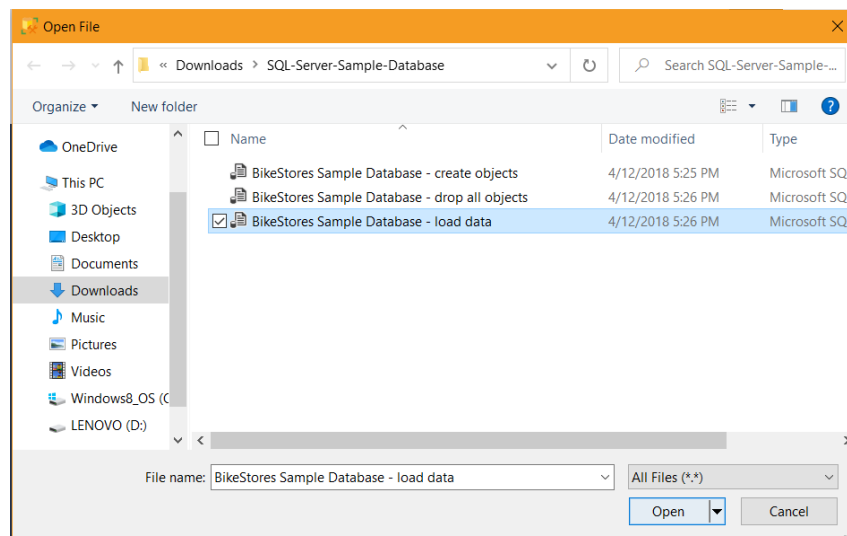
Select the **BikeStores Sample Database – create** objects.sql file and click the Open button



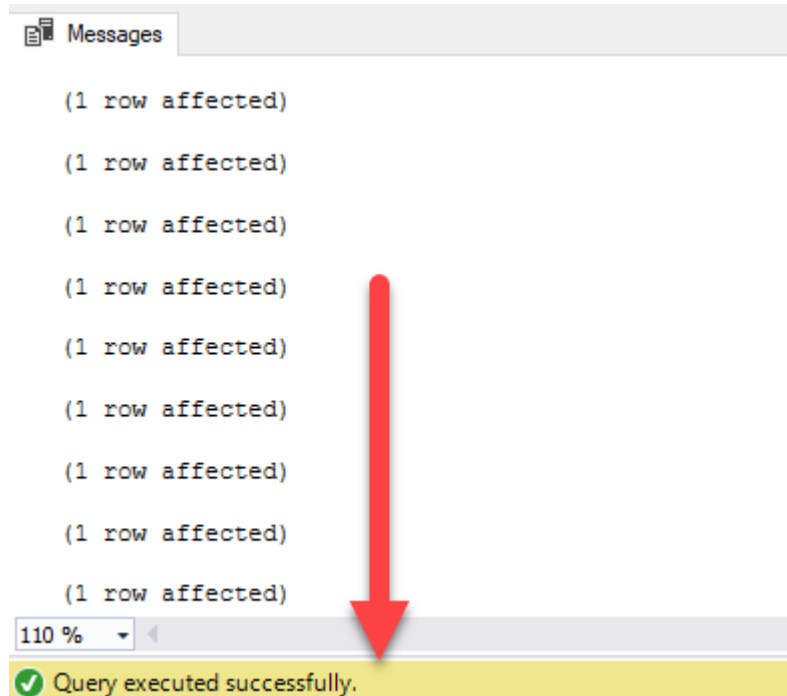
Click the **Execute** button to execute the SQL script



Choose the **BikeStores Sample Database – load data.sql** file and click the Open button.



Click the **Execute** button to load data into the tables.



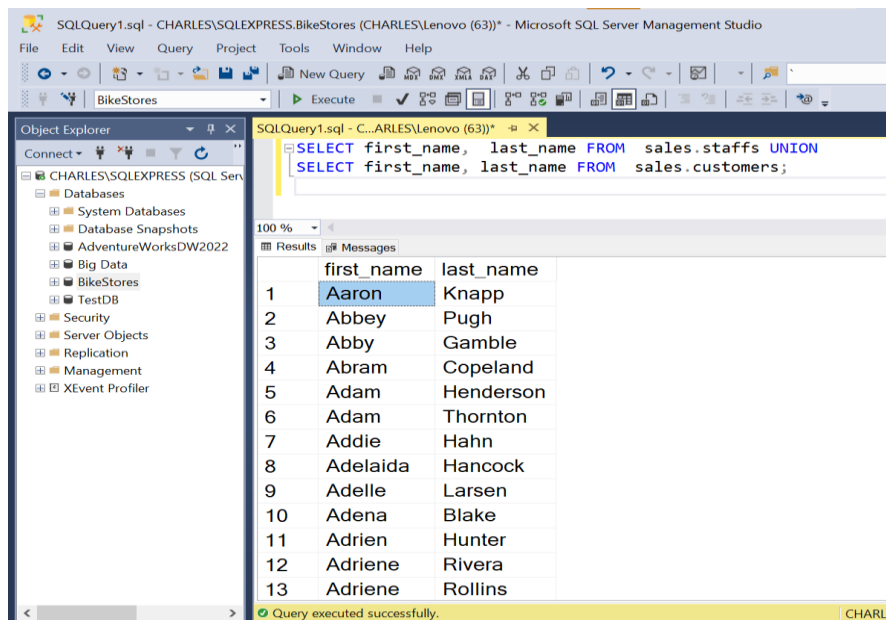
Execute query to display the values from production.products table.

QUERY OPERATIONS ON THE DATA WAREHOUSE

UNION OPERATION

SELECT first_name, last_name FROM sales.staffs UNION

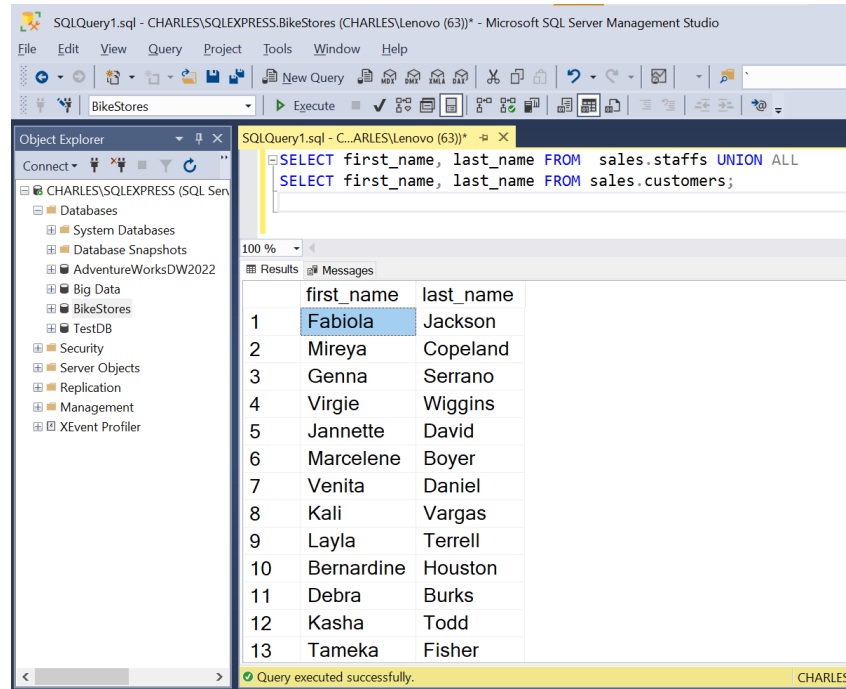
SELECT first_name, last_name FROM sales.customers;



UNION ALL OPERATION

SELECT first_name, last_name FROM sales.staffs UNION ALL

SELECT first_name, last_name FROM sales.customers;



The screenshot shows the Microsoft SQL Server Management Studio interface. The query editor contains the following SQL code:

```
SELECT first_name, last_name FROM sales.staffs UNION ALL  
SELECT first_name, last_name FROM sales.customers;
```

The query results are displayed in a table with two columns: first_name and last_name. The results are as follows:

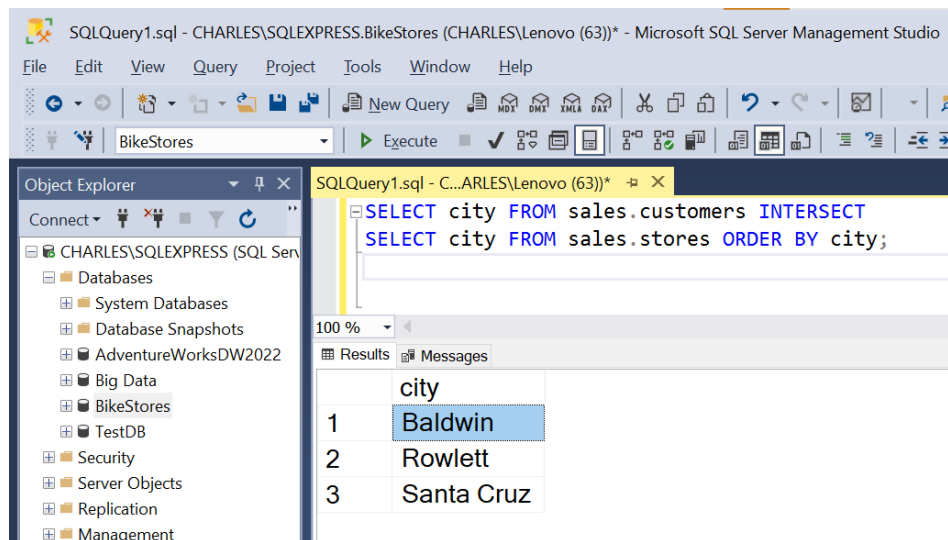
	first_name	last_name
1	Fabiola	Jackson
2	Mireya	Copeland
3	Genna	Serrano
4	Virgie	Wiggins
5	Jannette	David
6	Marcelene	Boyer
7	Venita	Daniel
8	Kali	Vargas
9	Layla	Terrell
10	Bernardine	Houston
11	Debra	Burks
12	Kasha	Todd
13	Tameka	Fisher

The status bar at the bottom indicates "Query executed successfully."

INTERSECT OPERATION

SELECT city FROM sales.customers INTERSECT

SELECT city FROM sales.stores ORDER BY city;



The screenshot shows the Microsoft SQL Server Management Studio interface. The query editor contains the following SQL code:

```
SELECT city FROM sales.customers INTERSECT  
SELECT city FROM sales.stores ORDER BY city;
```

The query results are displayed in a table with one column: city. The results are as follows:

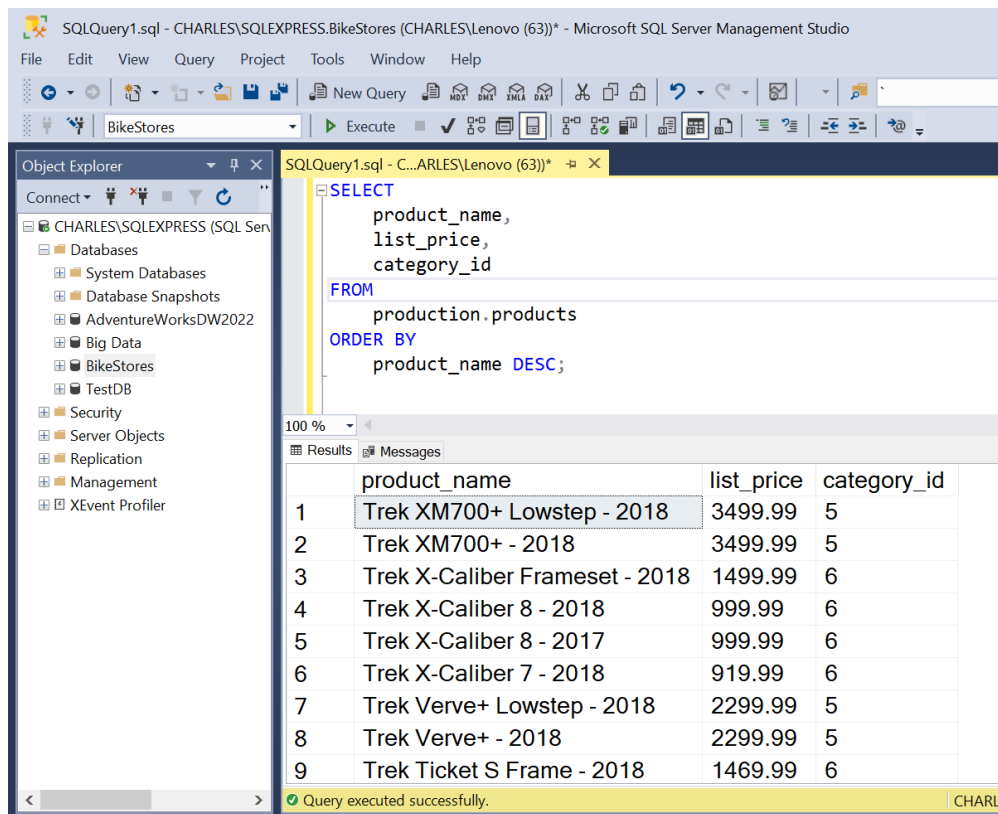
	city
1	Baldwin
2	Rowlett
3	Santa Cruz

The status bar at the bottom indicates "Query executed successfully."

JOIN OPERATION

INNER JOIN

```
SELECT
    product_name,
    list_price,
    category_id
FROM
    production.products
ORDER BY
    product_name DESC;
```



LEFT JOIN

```
SELECT
    product_name,
    order_id
FROM
    production.products p
LEFT JOIN sales.order_items o ON o.product_id = p.product_id
ORDER BY
    order_id;
```

SQLQuery1.sql - CHARLES\SQLEXPRESS.BikeStores (CHARLES\Lenovo (63)) - Microsoft SQL Server Management Studio

File Edit View Query Project Tools Window Help

Object Explorer

- CHARLES\SQLEXPRESS (SQL Ser...
- Databases
 - System Databases
 - Database Snapshots
 - AdventureWorksDW2022
 - Big Data
 - BikeStores
 - TestDB
- Security
- Server Objects
- Replication
- Management
- XEvent Profiler

SQLQuery1.sql - C:\ARLES\Lenovo (63))

```
SELECT
    product_name,
    order_id
FROM
    production.products p
LEFT JOIN sales.order_items o ON o.product_id = p.product_id
ORDER BY
    order_id;
```

100 %

Results Messages

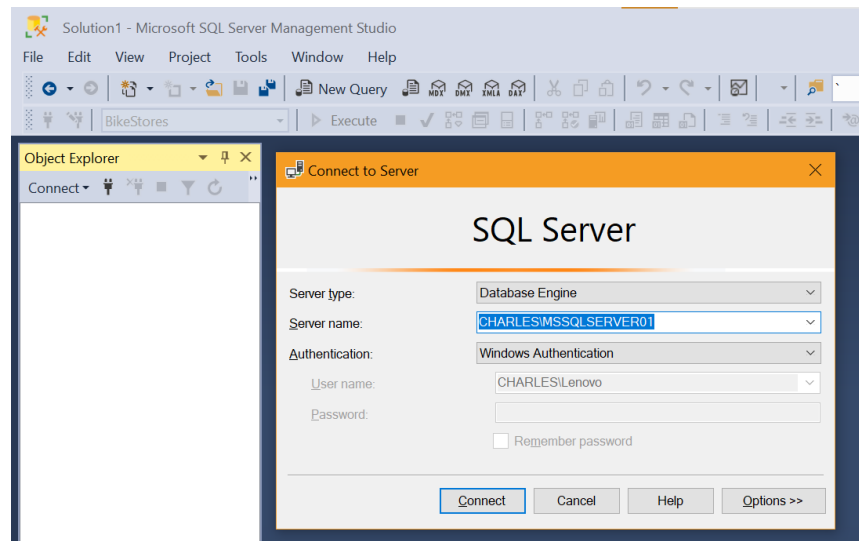
	product_name	order_id
1	Electra Savannah 1 (20-inch) - Girl's - 2018	NULL
2	Electra Townie Go! 8i Ladies' - 2018	NULL
3	Trek Checkpoint ALR 5 Women's - 2019	NULL
4	Trek Checkpoint ALR Frameset - 2019	NULL
5	Trek Precaliber 12 Girl's - 2018	NULL
6	Surly Krampus Frameset - 2018	NULL
7	Trek Checkpoint SL 5 Women's - 2019	NULL
8	Trek 820 - 2016	NULL
9	Trek Checkpoint ALR 4 Women's - 2019	NULL

Query executed successfully. CHARL

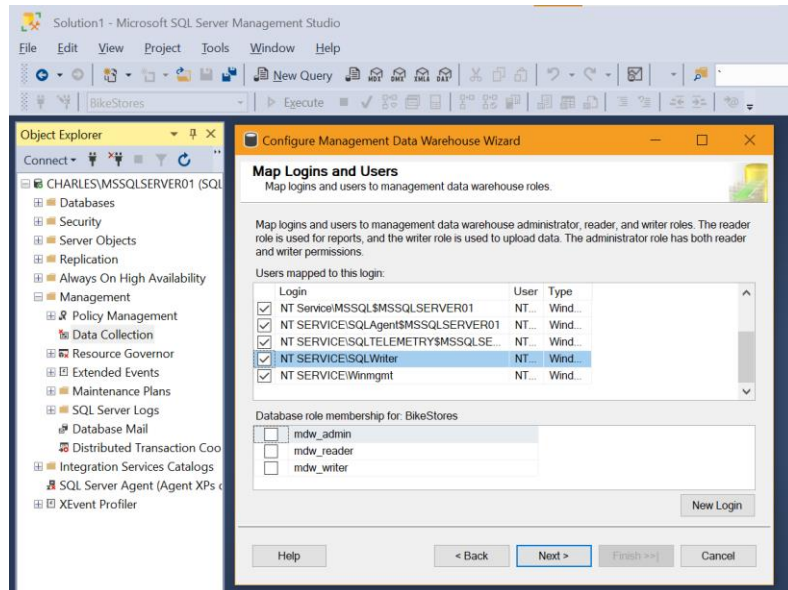
TASK 2b:

PROCEDURE:

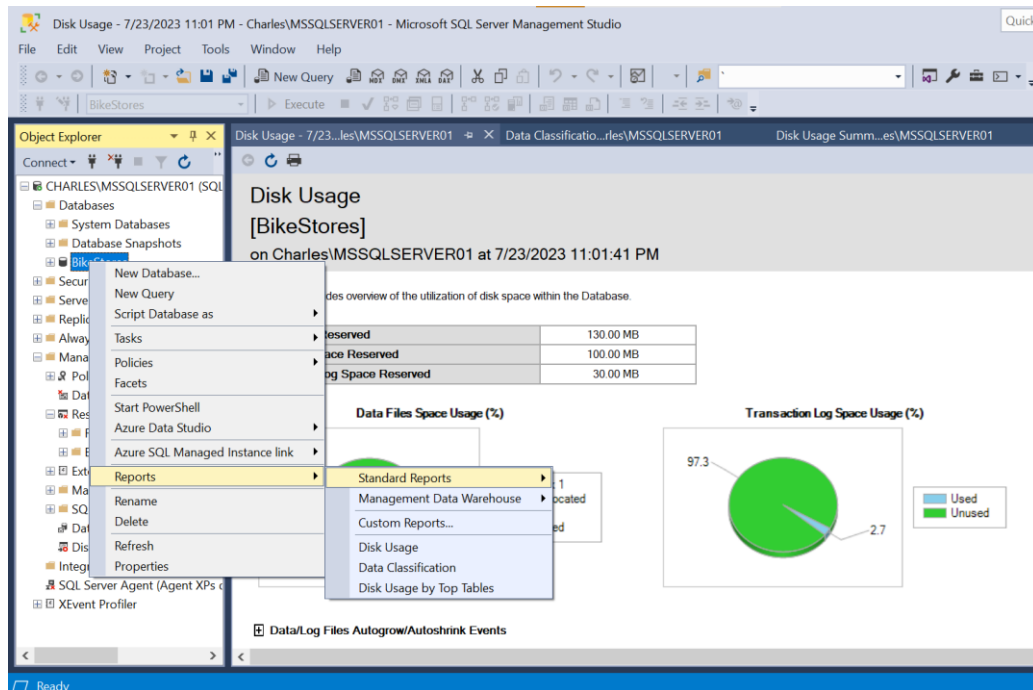
1. Install Microsoft SQL server Management Studio
2. Install MS Sql Server Developer Edition
3. Connect to the SQL Server



4. Choose Management -> Data Collection in the left side navigation bar
5. Right click on the Data collection - > Tasks -> Configure Management Data Warehouse and Configure the data warehouse settings
6. Select Server Name and Database name and click next button
7. Change the settings in map login and users
 - a. User mapped to this login
 - b. Database membership for BikeStores



8. Check the Settings and Click finish button
9. Monitor and generate the various reports for the BikeStores, Right click the Bikestores-> Reports -> Standard Reports and select the reports needed.



RESULT:

Date:

Task No: 3

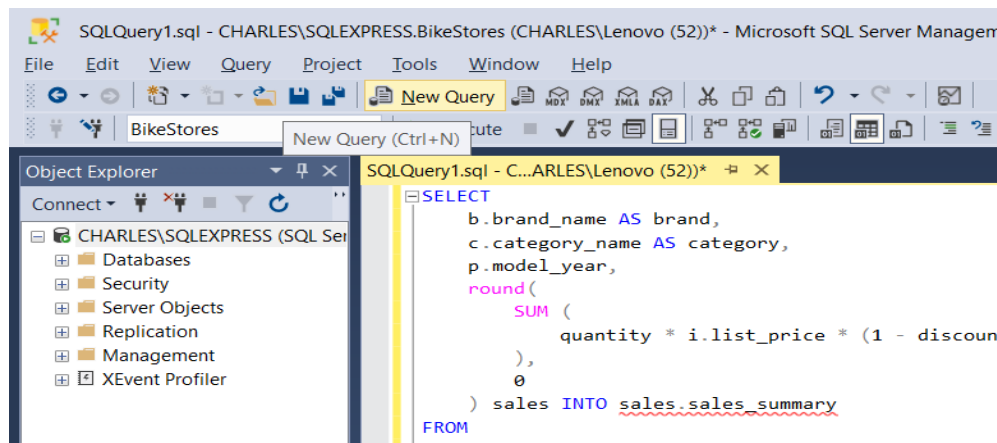
**Perform Data Cube Operations (OLAP Operations) using SQL Queries Rollup, Rolldown,
Slicing, Dicing
Database: MySQL**

AIM:

PROCEDURE:

1) create the sales.sales_summary table by using the following query:

```
SELECT
    b.brand_name AS brand,
    c.category_name AS category,
    p.model_year,
    round(
        SUM (
            quantity * i.list_price * (1 - discount)
        ),
        0
    ) sales INTO sales.sales_summary
FROM
    sales.order_items i
INNER JOIN production.products p ON p.product_id = i.product_id
INNER JOIN production.brands b ON b.brand_id = p.brand_id
INNER JOIN production.categories c ON c.category_id = p.category_id
GROUP BY
    b.brand_name,
    c.category_name,
    p.model_year
ORDER BY
    b.brand_name,
    c.category_name,
    p.model_year;
```



2) CUBE:

Refer task 2a:

The CUBE is a subclause of the GROUP BY clause that allows you to generate multiple grouping sets. The following illustrates the general syntax of the CUBE:

```
SELECT
    d1,
    d2,
    d3,
    aggregate_function (c4)
FROM
    table_name
GROUP BY
    d1,
    CUBE (d2, d3);
```

Example:

```
SELECT
    brand,
    category,
    SUM (sales) sales
FROM
    sales.sales_summary
GROUP BY
    CUBE (brand, category);
```

3) ROLLUP

Refer task 2a:

SQL Server ROLLUP syntax

The general syntax of the SQL Server ROLLUP is as follows:

```
SELECT
    d1,
    d2,
    d3,
    aggregate_function(c4)
FROM
    table_name
GROUP BY
    ROLLUP (d1, d2, d3);
```

Upload and configure the sample data warehouse (refer task2a)

Select the New Query, then new query tab will be opened.

ROLLUP Query

```
SELECT
    brand,
    category,
    SUM (sales) sales
FROM
    sales.sales_summary
GROUP BY
    ROLLUP (brand, category);
```

4) DRILLDOWN

This is a reverse of the ROLL UP operation. The data is aggregated from a higher level summary to a lower level summary/detailed data.

```
SELECT
    brand,
    category,
    SUM (sales) sales
FROM
    sales.sales_summary
GROUP BY
    ROLLUP (brand, category)
HAVING SUM (sales) > 20000;
```

5) SLICING:

A slice in a multidimensional array is a column of data corresponding to a single value for one or more members of the dimension. It helps the user to visualize and gather the information specific to a dimension.

```
SELECT
  *
FROM
  sales.customers
WHERE
  state = 'CA';
```

6) DICING:

Dicing is similar to slicing, but it works a little bit differently. Dicing, on the other hand, is more of a zoom feature that selects a subset over all the dimensions, but for specific values of the dimension.

```
SELECT
  city,
  COUNT (*)
FROM
  sales.customers
WHERE
  state = 'CA'
GROUP BY
  city
HAVING
  COUNT (*) > 10
ORDER BY
  city;
```

7) PIVOT

The steps to make a query a pivot table:

- 1) First, select a base dataset for pivoting.
- 2) Second, create a temporary result by using a derived table or common table expression (CTE)
- 3) Third, apply the PIVOT operator.

- First, select category name and product id from the production.products and production.categories tables as the base data for pivoting:

```
SELECT
  category_name,
  product_id
FROM
  production.products p
  INNER JOIN production.categories c
    ON c.category_id = p.category_id
```

- Second, create a temporary result set using a derived table:

```
SELECT * FROM (  
  SELECT  
    category_name,  
    product_id  
  FROM  
    production.products p  
    INNER JOIN production.categories c  
      ON c.category_id = p.category_id  
) t
```

➤ Third, apply the PIVOT operator:

```
SELECT * FROM  
(  
  SELECT  
    category_name,  
    product_id  
  FROM  
    production.products p  
    INNER JOIN production.categories c  
      ON c.category_id = p.category_id  
) t  
PIVOT(  
  COUNT(product_id)  
  FOR category_name IN (  
    [Children Bicycles],  
    [Comfort Bicycles],  
    [Cruisers Bicycles],  
    [Cyclocross Bicycles],  
    [Electric Bikes],  
    [Mountain Bikes],  
    [Road Bikes])  
) AS pivot_table;
```

RESULT:

Date:

Task No: 4

Implement matrix multiplication with Map Reduce

AIM:

Task 4.1

VMWare AND Apache Hadoop

To Install VMWare AND Apache Hadoop, follow the steps

STEP 1. First of all, enter to the official site of VMware and download VMware Workstation
<https://www.vmware.com/tryvmware/?p=workstation-w>

STEP 2. After downloading VMware workstation, install it on your PC

Step3: Installing Java 8 version.

Openjdk version "1.8.0_91"

OpenJDK Runtime Environment (build 1.8.0_91-8u91-b14-3ubuntu1~16.04.1-b14)

OpenJDK 64-Bit Server VM (build 25.91-b14, mixed mode)

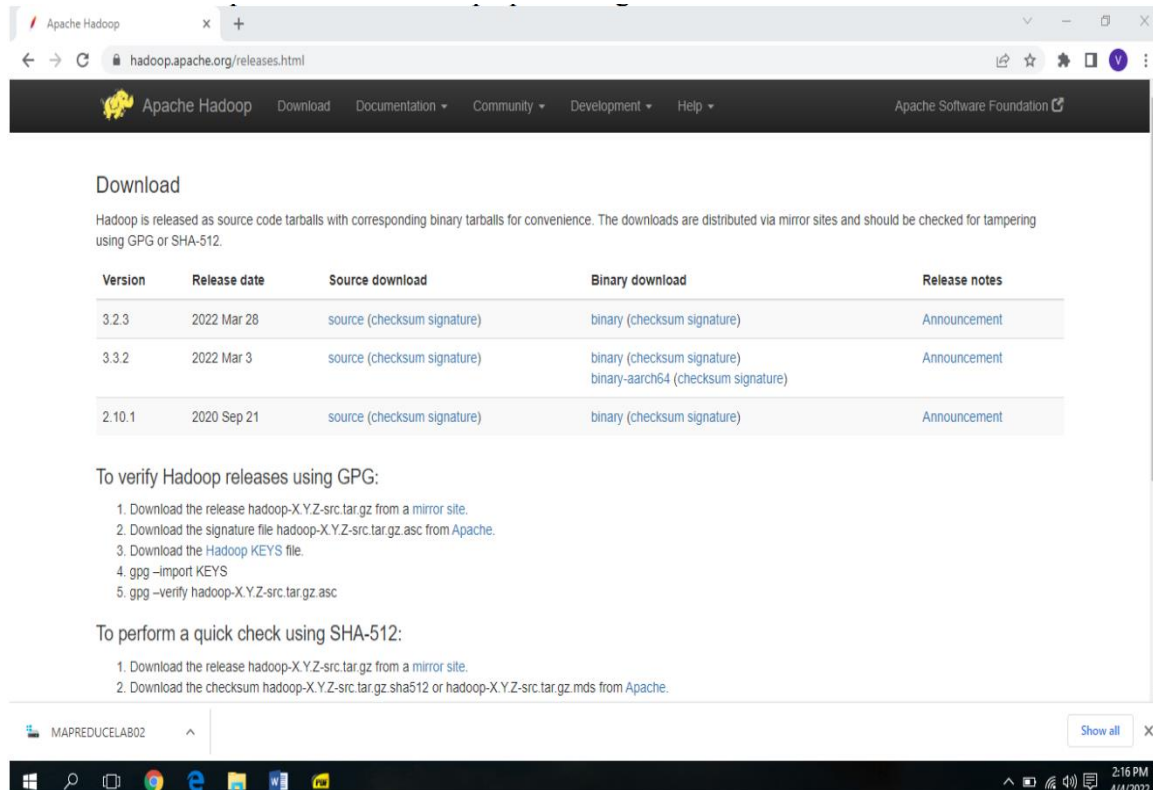
This output verifies that OpenJDK has been successfully installed.

Note: To set the path for environment variables. i.e. JAVA_HOME

Step4: Installing Hadoop

With Java in place, we'll visit the Apache Hadoop Releases page to find the most recent stable release. Follow the binary for the current release:

Download Hadoop from www.hadoop.apache.org



Task 4.2

Develop a MapReduce program to implement Matrix Multiplication.

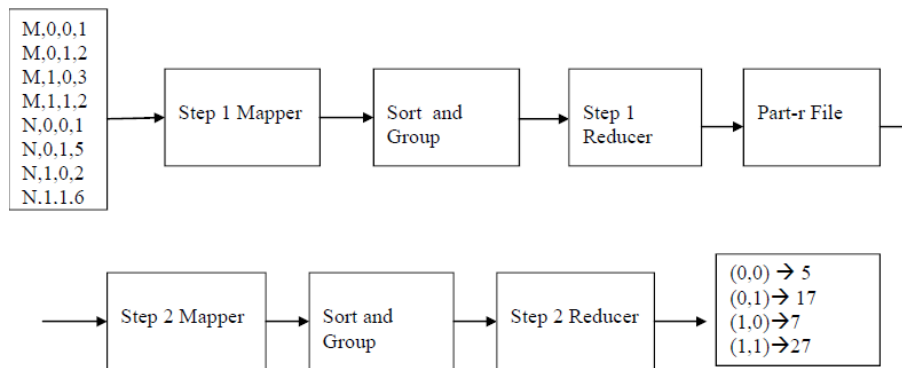
Assumes that you are using the Hadoop MapReduce framework and have two input matrices, Matrix A and Matrix B, stored in separate files.

Mapper Function: The mapper function takes as input a row from Matrix A and a column from Matrix B. It emits key-value pairs, where the key is the index of the resulting cell in the output matrix, and the value is the partial product of the corresponding elements from Matrix A and Matrix B.

Reducer Function:

The reducer function takes as input a key representing the index of a cell in the output matrix and a list of values representing the partial products for that cell. It calculates the final product by summing up the partial products and emits the result as the output.

Flow Diagram



The implementation should perform matrix multiplication using the MapReduce paradigm by distributing the computation across multiple mappers and reducers, resulting in efficient processing of large-scale matrices.

Task 4.3

Develop a MapReduce program to find the maximum temperature in each year.

A MapReduce program to find the maximum temperature in each year using the MapReduce paradigm. Assume that you are using the Hadoop MapReduce framework and have input data in the format: <Year> <Temperature>.

Mapper Function:

The mapper function takes as input a line from the input data file. It extracts the year and temperature from each line and emits key-value pairs, where the key is the year, and the value is the temperature.

Reducer Function:

The reducer function takes as input a year and a list of temperatures associated with that year. It iterates through the list of temperatures to find the maximum temperature and emits the result as the output.

RESULT:

Date:

Task No: 5

Write a Spark application to perform word count in the input file.

AIM:

5.1: Apache Spark Installation

PROCEDURE:

1. Download and install Java Development Kit (JDK) version 8 or higher, and ensure the RAM size, least 8 GB.
2. Download and install Python latest version from <https://www.python.org/>
3. Visit the Apache Spark website at <https://spark.apache.org/downloads.html> to download the latest stable release of Spark. [figure 1]
4. Type the following command in the command prompt to check the java and python version:
 - java --version
 - python --version
5. Create a new folder named Spark in the root of your C: drive and locate the Spark file you downloaded.
6. Right-click the file and extract it to C:\Spark using the tool you have on your system
7. Download the winutils.exe file for the underlying Hadoop version for the Spark installation, from the URL <https://github.com/cdarlint/winutils> [figure 2]
8. Configure Environment Variables
 - Click Start and type environment.
 - Select the result labeled Edit the system environment variables.
 - A System Properties dialog box appears. In the lower-right corner, click Environment Variables and then click New in the next window.
 - For Variable Name type SPARK_HOME. [figure 3,4]
 - For Variable Value type C:\Spark\spark-2.4.5-bin-hadoop2.7 and click OK. If you changed the folder path, use that one instead.
 - In the top box, click the Path entry, then click Edit. Be careful with editing the system path. Avoid deleting any entries already on the list.
 - The system highlights a new line. Enter the path to the Spark folder C:\Spark\spark-2.4.5-bin-hadoop2.7\bin. We recommend using %SPARK_HOME%\bin to avoid possible issues with the path. [figure 5,6,7]
 - Repeat this process for Hadoop and Java. For Hadoop, the variable name is HADOOP_HOME and for the value use the path of the folder you created earlier:

C:\hadoop. Add C:\hadoop\bin to the Path variable field, but we recommend using %HADOOP_HOME%\bin.

- For Java, the variable name is JAVA_HOME and for the value use the path to your Java JDK directory (in our case it's C:\Program Files\Java\jdk1.8.0_251).

9. Launch Spark, To start Spark, enter the command

- C:\Spark\spark-2.4.5-bin-hadoop2.7\bin>spark-shell
- If you set the environment path correctly, you can type spark-shell to launch Spark.
- Finally, the Spark logo appears, and the prompt displays the Scala shell.
- Open a web browser and navigate to <http://localhost:4040/>.

10. Launch pyspark, enter the command

```
C:\Spark\spark-2.4.5-bin-hadoop2.7\bin>pyspark
```

APACHE

Spark™





Lightning-fast unified analytics engine

DownloadLibraries ▾Documentation ▾ExamplesCommunity ▾Developers ▾

Download Apache Spark™

1. Choose a Spark release: 2.4.5 (Feb 05 2020)
2. Choose a package type: Pre-built for Apache Hadoop 2.7
3. Download Spark: [spark-2.4.5-bin-hadoop2.7.tgz](#)
4. Verify this release using the 2.4.5 [signatures](#), [checksums](#) and [project release KEYS](#).

Note that, Spark is pre-built with Scala 2.11 except version 2.4.2, which is pre-built with Scala 2.12.

 mapred	some binaries from 273 to 311
 mapred.cmd	some binaries from 273 to 311
 rcc	some binaries from 273 to 311
 winutils.exe	fixed exe and lib 265-312
 winutils.pdb	fixed exe and lib 265-312
 yarn	some binaries from 273 to 311
 yarn.cmd	some binaries from 273 to 311

Environment Variables

User variables for Goran

Variable	Value
OneDrive	C:\Users\Goran\OneDrive
Path	C:\Program Files\Intel\WiFi\bin\;C:\Program Files\Common Files\...
TEMP	C:\Users\Goran\AppData\Local\Temp
TMP	C:\Users\Goran\AppData\Local\Temp

1

New...

Edit...

Delete

Value

C:\WINDOWS\system32\cmd.exe

C:\Windows\System32\Drivers\DriverData

4

Windows_NT

C:\Program Files (x86)\Common Files\Oracle\Java\javapath;C:\WIN...

.COM;.EXE;.BAT;.CMD;.VBS;.VBE;.JS;.JSE;.WSF;.WSH;.MSC

AMD64

2

New...

Edit...

Delete

OK

Cancel

System Properties

Computer Name Hardware Advanced System Protection Remote

You must be logged on as an Administrator to make most of these changes.

Performance

Visual effects, processor scheduling, memory usage, and virtual memory

Settings...

User Profiles

Desktop settings related to your sign-in

Settings...

Startup and Recovery

System startup, system failure, and debugging information

Settings...


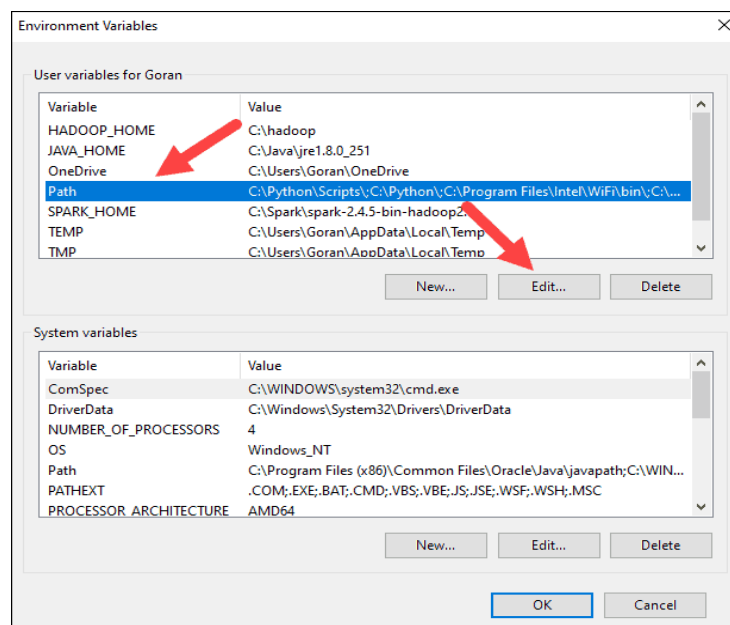
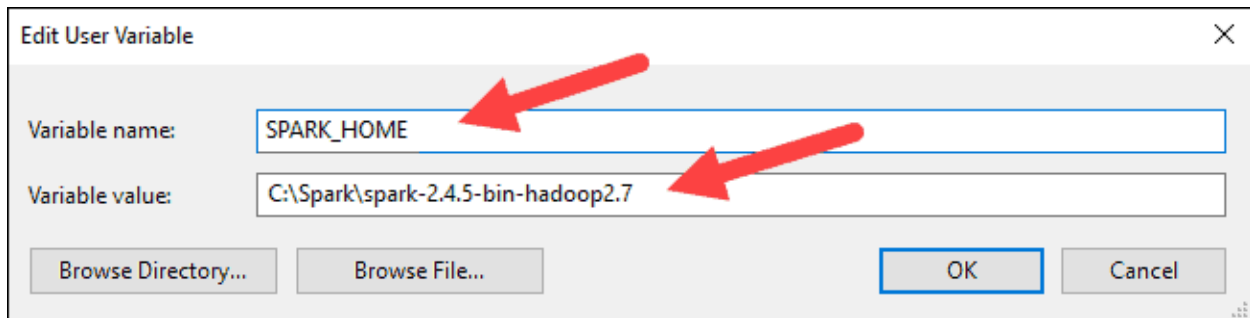
1

Environment Variables...

OK

Cancel

Apply



5.2: Spark application to perform word count

PROCEDURE:

1. Check the spark installation, environmental variables setup
2. Import necessary libraries from pyspark
3. Create a SparkConf and SparkContext (or SparkSession):
4. Load the input data
5. Read the input file and Calculating words count
6. Perform the word count operation
7. Save the output, Stop the SparkContext or SparkSession
8. Stopping Spark-Session and Spark context

Program:

```
import findspark

findspark.init()

from pyspark.sql import SparkSession

spark = SparkSession.builder\
    .master("local")\
    .appName('Firstprogram')\
    .getOrCreate()

sc=spark.sparkContext

text_file = sc.textFile("firstprogram.txt")

counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda x, y: x + y)

output = counts.collect()

for (word, count) in output:
    print("%s: %i" % (word, count))

sc.stop()

spark.stop()
```

File Name: firstprogram.txt

Chennai formerly known as Madras, is the capital city of Tamil Nadu, the southernmost Indian state.

Output:

Chennai: 1

Formerly: 1

Known: 1

as: 1

Madras: 1

Is: 1

the: 2

Capital: 1

City: 1

Of: 1

Tamil: 1

Nadu: 1

southernmost: 1

Indian: 1

state.: 1

RESULT:

Date:

Task No: 6
Implement CRUD operations on Casandra

AIM:

PROCEDURE:

1. Install Java Development Kit (JDK) version 8 or higher.
2. Install Python stable version
3. Visit the Apache Cassandra website at <https://cassandra.apache.org/download/> to download the latest stable release of Cassandra.
4. Set up the required, environment variables.
5. Modify the configuration files cassandra.yaml to adjust the settings according to your requirements.
6. Open the command prompt at bin folder and Start the Cassandra server using the command “cassandra”
7. Open the another command prompt at bin folder and start the command-line interface (CLI) using the command “cqls”
8. Create the keyspace and tables, then execute the CRUD operations – create, update, read, delete.
9. Execute various no-sql queries and find the results.
10. Stop the cassandra server.

Create Keyspace:

A keyspace is like RDBMS database which contains column families, indexes, user defined types, data center awareness, strategy used in keyspace, replication factor, etc.

CREATE KEYSPACE <identifier> WITH <properties>

Or

Create keyspace KeyspaceName with replicaton={'class':strategy name,
'replication_factor': No of replications on different nodes}

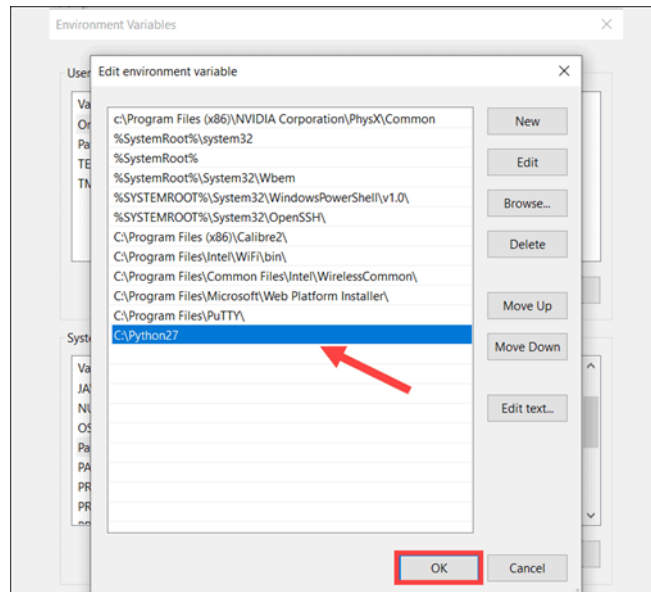

Alter Keyspace:

ALTER KEYSPACE <identifier> WITH <properties>

Or

ALTER KEYSPACE "KeySpace Name"

WITH replication = {'class': 'Strategy name', 'replication_factor' : 'No.Of replicas'};



Apache Software Foundation / Apache Cassandra / Download

 **CASSANDRA™**

Home Download Documentation Community Blog

Downloading Cassandra

Latest version

Download the latest Apache Cassandra 3.11 release: **3.11.6** (pgp, sha256 and sha512), released on 2020-02-14.

Older supported releases

The following older Cassandra releases are still supported:

- Apache Cassandra 3.0 is supported until **6 months after 4.0 release (date TBD)**. The latest release is 3.0.20 (pgp, sha256 and sha512), released on 2020-02-14.
- Apache Cassandra 2.2 is supported until **4.0 release (date TBD)**. The latest release is 2.2.16 (pgp, sha256 and sha512), released on 2020-02-14.
- Apache Cassandra 2.1 is supported until **4.0 release (date TBD)** with **critical fixes only**. The latest release is 2.1.21 (pgp, sha256 and sha512), released on 2019-02-11.

Older (unsupported) versions of Cassandra are [archived here](#).

Drop Keyspace:

DROP keyspace KeyspaceName ;

Create Table Operation:

Syntax of Create Operation-

CREATE (TABLE | COLUMNFAMILY) <tablename>

('<column-definition>' , '<column-definition>')

(WITH <option> AND <option>)

Sample No-SQL queries

```
CREATE TABLE student(  
    student_id int PRIMARY KEY,  
    student_name text,  
    student_city text,  
    student_fees varint,  
    student_phone varint  
);
```

CURD OPERATION:**Create Data:**

```
INSERT INTO student (student_id, student_fees, student_name)
```

```
VALUES(1,5000, 'Ajeet');
```

```
INSERT INTO student (student_id, student_fees, student_name)
```

```
VALUES(2,3000, 'Kanchan');
```

```
INSERT INTO student (student_id, student_fees, student_name)
```

```
VALUES(3, 2000, 'Shivani');
```

READ Data

```
SELECT * FROM student;
```

```
SELECT * FROM student WHERE student_id=2;
```

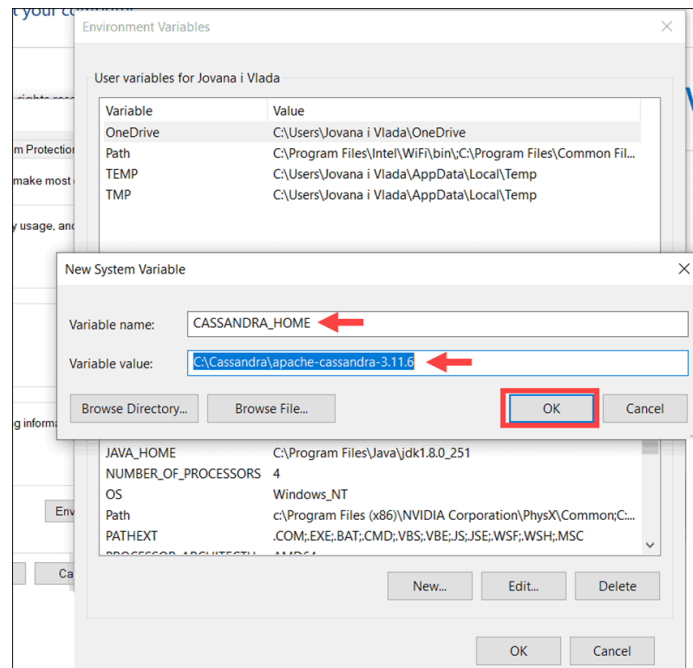
Update Data:

```
UPDATE student SET student_fees=10000,student_name='Rahul'
```

```
WHERE student_id=2;
```

DELETE Data:

```
DELETE student_fees FROM student WHERE student_id=4;
```



```
C:\Windows\System32\cmd.exe
Microsoft Windows [Version 10.0.19045.3324]
(c) Microsoft Corporation. All rights reserved.

C:\cassandra\apache-cassandra-3.0.29\bin>cassandra
WARNING! Powershell script execution unavailable.
Please use 'powershell Set-ExecutionPolicy Unrestricted'
on this user-account to run cassandra with fully featured
functionality on this platform.
Starting with legacy startup options
Starting Cassandra Server
```

```
C:\Windows\System32\cmd.exe
Microsoft Windows [Version 10.0.19045.3324]
(c) Microsoft Corporation. All rights reserved.

C:\cassandra\apache-cassandra-3.0.29\bin>cqlsh
File "C:\cassandra\apache-cassandra-3.0.29\bin\cqlsh.py", 1
```

RESULT:

Date:

Task No: 7

Implementing Producer and Consumer problem in kafka.

AIM:

PROCEDURE:

1. Install Java Development Kit (JDK) version 8 or higher.
2. Install ZooKeeper: Kafka relies on ZooKeeper for distributed coordination. Kafka comes with an inbuilt Zookeeper
3. Visit the Apache Kafka website (<https://kafka.apache.org/downloads>) and download the latest version of Kafka for Windows.
4. Extract the contents of the Apache Kafka ZIP file to a required directory
5. Configure Environment Variables and path in System and User settings.
6. Start zookeeper by execute the zookeeper-server-start.bat script and pass the zookeeper configuration file path
`cd C:\bigdata\kafka2.5`
`start cmd /k bin\windows\zookeeper-server-start.bat config\zookeeper.properties`
7. Start Kafka by executing the kafka-server-start.bat script and pass broker configuration file path.
`cd C:\bigdata\kafka2.5`
`start cmd /k bin\windows\kafka-server-start.bat config\server.properties`
8. Open a new command prompt, and create new Kafka topic.
`bin\windows\kafka-topics.bat --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1 --topic test`
Created topic test.
9. List all the topics to verify the created topic is present in this Topics list
`bin\windows\kafka-topics.bat --list --bootstrap-server localhost:9092`
test
10. Produce some messages and submit to test topic. I added two messages i.e. "Hello" and "Kafka".
`bin\windows\kafka-console-producer.bat --bootstrap-server localhost:9092 --topic test`

```
Command Prompt - bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --...
C:\Users\Lokesh>cd E:\devsetup\bigdata\kafka2.5
C:\Users\Lokesh>e:
E:\devsetup\bigdata\kafka2.5>bin\windows\kafka-topics.bat --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1 --topic test
Created topic test.
E:\devsetup\bigdata\kafka2.5>bin\windows\kafka-topics.bat --list --bootstrap-server localhost:9092
test
E:\devsetup\bigdata\kafka2.5>bin\windows\kafka-console-producer.bat --bootstrap-server localhost:9092 --topic test
>Hello
>Kafka
>Terminate batch job (Y/N)? y
E:\devsetup\bigdata\kafka2.5>bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic test --from-beginning
Hello
Kafka
```

11. Consume the messages and submit to test topic.

```
bin\windows\kafka-console-consumer.bat --bootstrap-server localhost:9092 --topic test --from-beginning
```

Hello

Kafka

RESULT:

Date:

Task No: 8
Implement basis commands in HIVE.

AIM:

8 A)

PROCEDURE:

1. Install Java Development Kit (JDK) version 8 or higher.
2. Install Hadoop-2.9.2, or use any other STABLE version for Hadoop.
3. Install MySQL Query Browser
4. Download and install Hive-3.1.2, or use any other STABLE version for Hive.
5. Set the environmental variables for HIVE_HOME and Path variables %HIVE_HOME%\bin
6. Open a NEW Command Window, check the path settings
`echo %HIVE_HOME%`
7. Edit the Properties of hive-site.xml and Replace the value for <Your IP Address> with the IP Address of your System and replace <Your drive Folder> with the Hive folder Path.
8. Removing Special Characters and Adding few More Properties in hive-site.xml File.
9. Creating Hive User in MySQL Database for reading and writing data from it.
10. Grant permission to Users and Creating Metastore
11. Starting Hadoop, Hive Metastore and Hive shell prompt
`start-all.cmd`
`hive --service metastore`
12. Open a new cmd window and run the below command to start Hive
`> hive`

Hive Basic Commands:


Show Databases:

`SHOW DATABASES;`

Use Database:

`USE database_name;`

Show Tables:



SHOW TABLES;

Create a Database:

CREATE DATABASE new_database;

Create a Table:

CREATE TABLE my_table (

id INT,

name STRING,

age INT

);

Load Data into a Table:

LOAD DATA LOCAL INPATH '/path/to/datafile.csv' OVERWRITE INTO TABLE my_table;

Query Data:


SELECT * FROM my_table WHERE age > 25;

Insert Data into a Table:

INSERT INTO TABLE my_table VALUES (1, 'John', 30);

Describe Table:

DROP TABLE my_table;



8 B)

PROCEDURE:

1. Install Java Development Kit (JDK) version 8 or higher.
2. Install Hadoop-2.9.2, or use any other STABLE version for Hadoop.
3. Install MySQL Query Browser and configure the Hadoop, msq and Hive
4. Starting Hadoop, Hive Metastore and Hive shell prompt
start-all.cmd
hive --service metastore
5. Open a new cmd window and run the below command to start Hive
> hive
6. Execute the various Hql Commands in Hive shell
7. Stop the hive and other server

HQL COMMANDS:

Create Database

```
hive> create database demo;
```

```
hive> create a database if not exists demo;
```

```
hive> show databases;
```

```
hive> describe database extended demo;
```

Create Table

```
hive> create table demo.employee (Id int, Name string , Salary float)
```

```
row format delimited
```

```
fields terminated by ',' ;
```

```
hive> create table if not exists demo.employee (Id int, Name string , Salary float)
```

```
row format delimited
```

```
fields terminated by ',' ;
```

Alter Table

```
hive> alter table old_table_name rename to new_table_name;
```

```
hive> Alter table emp rename to employee_data;
```

Drop Table

```
hive> drop table new_employee;
```

Aggregate Functions in Hive

employee_data

Id	Name	Salary
1	Gaurav	30000
2	Aryan	20000
3	Vishal	40000
4	John	10000
5	Henry	25000
6	William	9000
7	Lisa	25000
8	Ronit	20000

Load the below data file “empl_details” from local drive to HQL table “employee_data” and perform HQL function.

```
hive> create table employee_data (Id int, Name string , Salary float)
```

```
row format delimited
```

```
fields terminated by ',' ;
```

```
hive> load data local inpath '/home/codegyani/hive/emp_details' into table employee_data;
```

```
hive> select max(Salary) from employee_data;
```

GROUP BY and HAVING Clause

```
hive> select department, sum(salary) from emp group by department;
```

```
hive> select department, sum(salary) from emp group by department having sum(salary)>=35000;
```

View and Indexes

HQL Views:

Views are generated based on user requirements and save any result set data as a view.

```
hive> CREATE VIEW emp_30000 AS
```

```
SELECT * FROM employee
```

```
WHERE salary>30000;
```

Dropping a View

```
hive> DROP VIEW emp_30000;
```

Creating an Index:

An Index is nothing but a pointer on a particular column of a table. Creating an index means creating a pointer on a particular column of a table.

```
hive> CREATE INDEX inedx_salary ON TABLE employee(salary)
```

```
AS 'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler';
```

Dropping an Index:

```
hive> DROP INDEX index_salary ON employee;
```

RESULTS:

Date:

Task No: 9

Write Pig Latin scripts sort, group, join, project, and filter the data

AIM:

PROCEDURE:

- 1) Install the prerequisite Stable software's Java Development Kit and Java Runtime Environment, Apache Hadoop.
- 2) Visit the Apache Pig download page: <https://pig.apache.org/downloads.html>
- 3) Download the latest stable release of Pig.
- 4) Extract the Pig Archive and Set Environment Variables for java, Apache Hadoop and Apache Pig

For Java:

JAVA_HOME=C:\Program Files\Java\jdk-1.8

Path = C:\Program Files\Java\jdk-1.8\bin

For Hadoop:

HADOOP_HOME=C:\ApacheHadoop2.9.2

Path = C:\ApacheHadoop2.9.2\bin

For Pig:

PIG_HOME = C:\Apachepig

Path = C:\Apachepig\bin

Path = C:\Apachepig\conf

- 5) Verify the Paths, Run following commands in a NEW Command Window
echo %PIG_HOME%
- 6) Open the pig. cmd file in edit mode, and change the value of the HADOOP_BIN_PATH
Old value:- %HADOOP_HOME%\bin
New Value:- %HADOOP_HOME%\libexec
- 7) Edit Pig Configuration, go to the conf directory within your Pig installation directory, rename the pig.properties.template file to pig.properties. and set the exectype property to "local"
exectype=local
- 8) Start Apache Pig, run the following command in a new Command Prompt as administrator
C:\Users\Lenovo>echo %PIG_HOME%
O/P: C:\ApachePig
C:\Users\Lenovo>pig -version

Output:

Pig Latin scripts:

Input.txt

Rajiv,42

siddarth,45

Rajesh,40

Preethi,23

Trupthi,34

Archana,21

Robin,22

BOB,23

Maya,23

Sara,25

David,23

Maggy,22

Addressfile.txt

Rajiv,Chennai

Rajesh,Delhi

Trupthi,Hyderabad

Robin,Pune

Maya,Hyderabad

Anderson,Chennai

Antolina,Chennai

Apache Pig version 0.17.0 (r1797386)

compiled Jun 02 2017, 15:41:58

C:\Users\Lenovo>pig

Grunt Shell started:

grunt>

Load Data:

```
data = LOAD './input.txt' USING PigStorage(',') AS (name:chararray, age:int);
```

Sort Operator:

```
sortbyage = ORDER data BY age ASC|DESC;  
dump sortbyage;
```

Group Operator:

```
grouped_data = GROUP data BY age;  
dump grouped_data;
```

Filter Operator:

```
filterbyage = FILTER data BY (age>40);  
dump filterbyage;
```

Inner Join Operator:

```
table1 = LOAD './input.txt' USING PigStorage(',') AS (name:chararray, age:int);  
table2 = LOAD './addressfile.txt' USING PigStorage(',') AS (name:chararray, address:chararray);
```

```
joinbyname = JOIN table1 BY name, table2 BY name;
```

OuterJoin Operator;

```
LO = JOIN table1 BY name LEFT OUTER, table2 BY name;  
dump LO;
```

```
RO = JOIN table1 BY name RIGHT, table2 BY name;  
dump RO;
```

```
FO = JOIN table1 BY name FULL OUTER, table2 BY name;  
dump FO;
```

Store Operator:

```
grouped_data = GROUP data BY age;  
result = FOREACH grouped_data GENERATE group AS age, COUNT(data) AS count;  
STORE result INTO 'output';
```

RESULT:

.

Date:

Task No: 10

- i.) Construct the Pig Latin Scripts to find character Count**
- ii. Construct the Pig Latin Scripts to find a max temp for each and every year**

AIM:

PROCEDURE:

Construct the Pig Latin Scripts to find character Count

- 1) Load the data from LOCAL file system
- 2) Convert the Sentence into words using TOKENIZE Function with delimiter
- 3) Convert Column into Rows
- 4) Apply GROUP BY and Generate word count
- 5) Dump the word count

Construct the Pig Latin Scripts to find a max temp for each and every year

- 1) Load your data (assuming it's in a tab-delimited file)
- 2) Extract the year from the date column (assuming date is in 'YYYY-MM-DD' format)
- 3) Group the data by year
- 4) Calculate the maximum temperature for each year
- 5) Store the result in an output file

Pig Latin Scripts:

- (i) Construct the Pig Latin Scripts to find character Count

File Name: wordcount.pig

```
data = LOAD './countwords.txt' USING PigStorage() AS (line:chararray);
```

```
Words = FOREACH data GENERATE FLATTEN(TOKENIZE(line,' ')) AS word;
```

```
Grouped = GROUP Words BY word;
```

```
wordcount = FOREACH Grouped GENERATE group, COUNT(Words);
```

```
DUMP wordcount;
```

- (ii) Construct the Pig Latin Scripts to find a max temp for each and every year

File Name: tempanalysis.pig

OUTPUT:

```
grunt>exec wordcount.pig
```

(a,2)

(is,2)

(This,1)

(post,1)

(hadoop,2)

(bigdata,1)

(technology,1)

```
grunt>exec tempanalysis.pig
```

(2020,38.5)

(2021,36.7)

(2022,40.8)

(2023,42.8)

```
raw_data = LOAD './temperature.txt' USING PigStorage('\t') AS (date:chararray, temperature:double);
data_with_year = FOREACH raw_data GENERATE SUBSTRING(date, 0, 4) AS year, temperature;
grouped_data = GROUP data_with_year BY year;

max_temp_per_year = FOREACH grouped_data GENERATE group AS year,
MAX(data_with_year.temperature) AS max_temperature;

STORE max_temp_per_year INTO 'your_output_file.txt' USING PigStorage('\t');
```

Input Files:

countwords.txt

This is a hadoop post
hadoop is a bigdata technology

temperature.txt

9/12/2020	33.4
7/13/2021	30.45
9/14/2020	28.6
8/15/2023	30.4
6/16/2022	33.5
9/17/2023	38.6
9/18/2020	38.5
5/19/2023	36.4
9/20/2023	35.7
9/21/2021	36.7
2/22/2023	39.4
9/23/2023	40.2
3/24/2022	40.6
1/25/2022	40.8
9/26/2023	42.8

RESULT:

Date:

Task No: 11

A . Collect Social Media Data from a Twitter to a Local File with the Topic 'covid 19'

AIM:

PROCEDURE:

1. Install tweepy and json module in pycharm
2. Import tweepy and json module.
3. Create an account in Twitter Developer mode.
4. Get the Twitter API credentials – consumer key, consumer secret, access token, access token secret.
5. Authenticate with Twitter API
6. Search query and set, number of tweets to collect
7. Create File to save the collected data with file name 'covid19_tweets.json'
8. Collect tweets and save to a file

PROGRAM:

```
import tweepy
import json

consumer_key = 'KbKxl06BkeqAcwNhJnBMN5s5M'
consumer_secret = 'LaVFLQhzy57mVDo9dbT60QOXnN6Ajm09Q6PAmSgPGoozZqhM9'
access_token = '1709726896194244609-w4KYxcGTyA4GHjlO7kJHzQlCWivVvSm'
access_token_secret = 'G8uyaHaIXoDUMFzVW3Ao6sAWRZ328Y2B6RraAeEU8p4e3'

auth = tweepy.OAuth2BearerHandler(consumer_key, consumer_secret)
auth.set_access_token(access_token, access_token_secret)
api = tweepy.API(auth, wait_on_rate_limit=True)
search_query = 'covid 19'
num_tweets = 1000
output_file = 'covid19_tweets.json'
tweets = []
for tweet in tweepy.Cursor(api.search, q=search_query, lang='en').items(num_tweets):
    tweets.append(tweet._json)

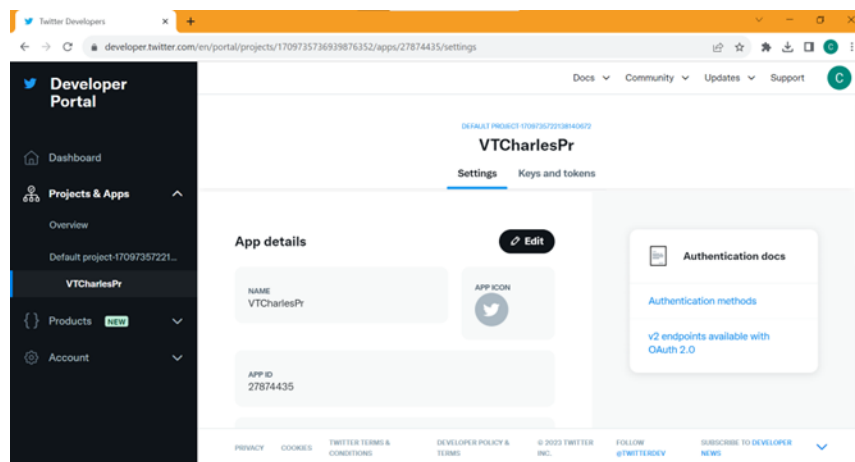
with open(output_file, 'w', encoding='utf-8') as file:
    json.dump(tweets, file, ensure_ascii=False, indent=4)

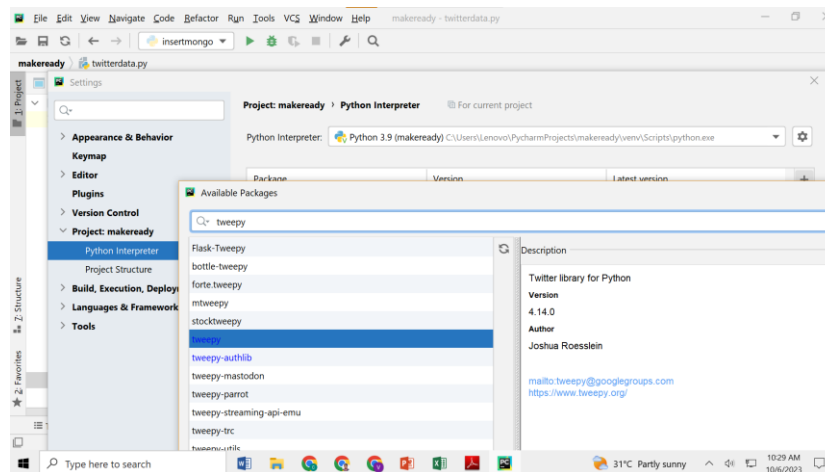
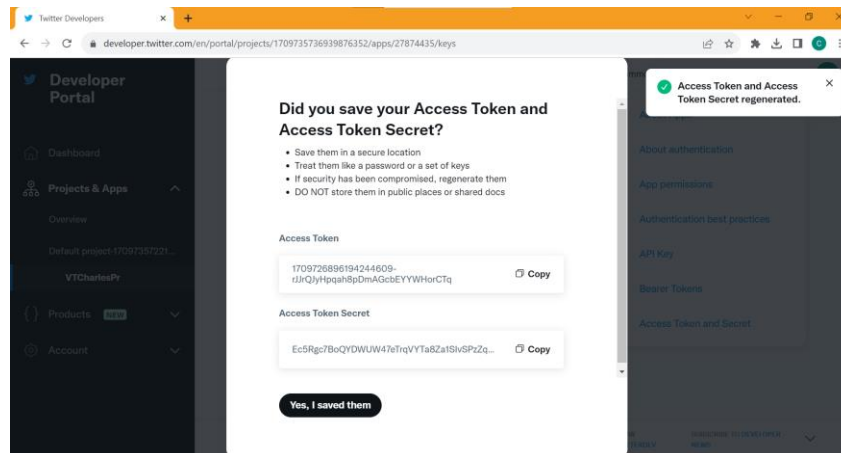
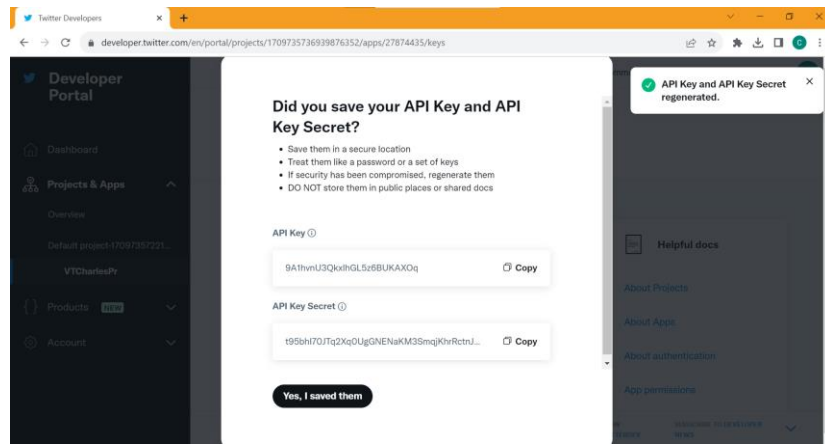
print(f'{len(tweets)} tweets collected and saved to {output_file}')
```


OUTPUT:

'covid19_tweets.json'

```
[{
  "State": "Andaman and Nicobar Islands",
  "Confirmed": 7651,
  "Recovered": 7518,
  "Deaths": 129,
  "Active": 4,
  "Last_Updated_Time": "13/08/2021 23:27:22",
  "Migrated_Other": 0,
  "State_code": "AN",
  "Delta_Confirmed": 0,
  "Delta_Recovered": 0,
  "Delta_Deaths": 0,
  "State_Notes": ""
},
{
  "State": "Andhra Pradesh",
  "Confirmed": 2066450,
  "Recovered": 2047722,
  "Deaths": 14373,
  "Active": 4355,
  "Last_Updated_Time": "13/08/2021 23:27:22",
  "Migrated_Other": 0,
  "State_code": "AP",
  "Delta_Confirmed": 0,
  "Delta_Recovered": 0,
  "Delta_Deaths": 0,
  "State_Notes": ""
}]
```







B. Download and Set Up MongoDB Server and a Client Mongoddb Compass.

PROCEDURE:

1. Installing MongoDB on Windows:
2. Go to the official MongoDB website (<https://www.mongodb.com/try/download/community>) and download the MongoDB Community Server for Windows.
3. Run the downloaded installer (.msi file).
4. Choose a custom installation if you want to specify a custom installation directory or options.
5. After installation, MongoDB should start automatically as a Windows service. Check this in the Windows Services Manager.
6. Open Mongoddb Compass, By default, MongoDB listens on port 27017.
7. Import the COIVD19, state-wise data into the Mongoddb

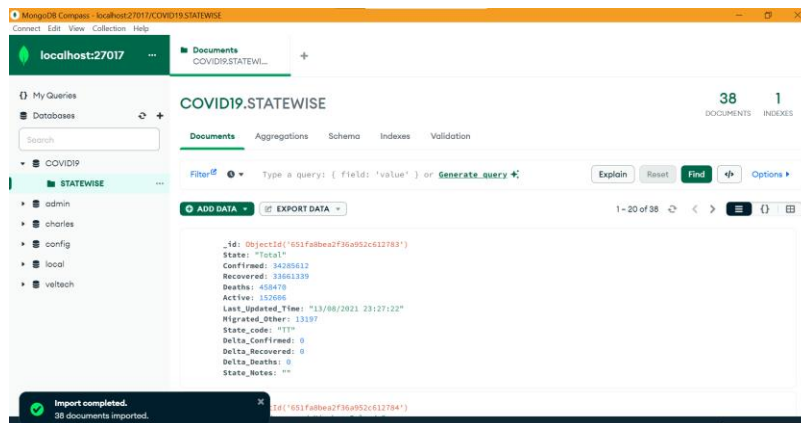
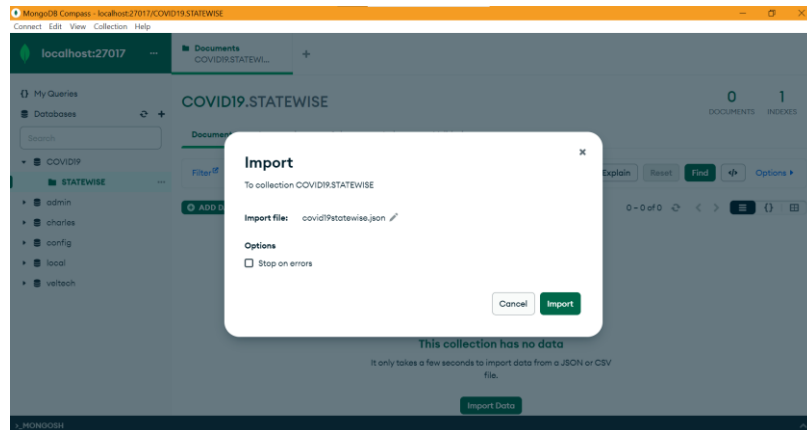
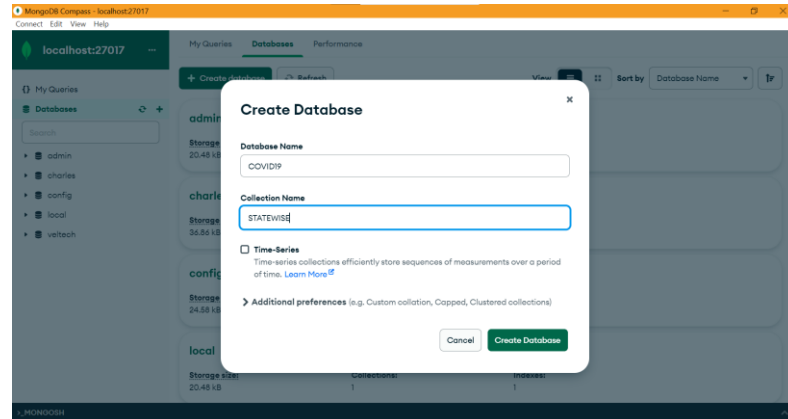
SAMPLE INPUT FILE: [File Name: covid19statewise.json]

```
[{
  "State": "Andaman and Nicobar Islands",
  "Confirmed": 7651,
  "Recovered": 7518,
  "Deaths": 129,
  "Active": 4,
  "Last_Updated_Time": "13/08/2021 23:27:22",
  "Migrated_Other": 0,
  "State_code": "AN",
  "Delta_Confirmed": 0,
  "Delta_Recovered": 0,
  "Delta_Deaths": 0,
  "State_Notes": ""
},
{
  "State": "Andhra Pradesh",
  "Confirmed": 2066450,
  "Recovered": 2047722,
  "Deaths": 14373,
  "Active": 4355,
  "Last_Updated_Time": "13/08/2021 23:27:22",
  "Migrated_Other": 0,
  "State_code": "AP",
  "Delta_Confirmed": 0,
  "Delta_Recovered": 0,
  "Delta_Deaths": 0,
  "State_Notes": ""
}]
```

OUTPUT:

Import Completed

38 Documents Imported



RESULT:

Date:

Task No: 12

- A. Retrieve Analytic Information given from MongoDB**
- i. For each “place_type”, Find total favorite_count**
 - ii. For each "country_code", find total "retweet_count"**

AIM:

.

PROCEDURE:

1. Install pymongo in Pycharm IDE
2. Import the pymongo module in the analytics program
3. Configure the mongodb server using mongodb compass
4. Upload the country dataset in to Mongodb
5. Get the MongoDB URI, Database name and Collection Name
6. Establish the MongoDB connection using MongoDB URI, Database name and Collection Name
7. Aggregation pipeline to calculate the total favorite_count, retweet_count
8. Generate the results using aggregate functions.
9. Print the results, total favorite_count, retweet_count

PROGRAM:

```
import pymongo
mongo_uri = "mongodb://localhost:27017/"
database_name = "Country"
collection_name = "Tweets"
client = pymongo.MongoClient(mongo_uri)
db = client[database_name]
collection = db[collection_name]

pipeline = [
    {
        "$group": {
            "_id": "$country_code",
            "total_favorite_count": {"$sum": "$favorite_count"},
            "total_retweet_count": {"$sum": "$retweet_count"}
        }
    }
]

result = list(collection.aggregate(pipeline))

if result:
    total_favorite_count = result[0]["total_favorite_count"]
```

OUTPUT:

Total_favorite_count: 3

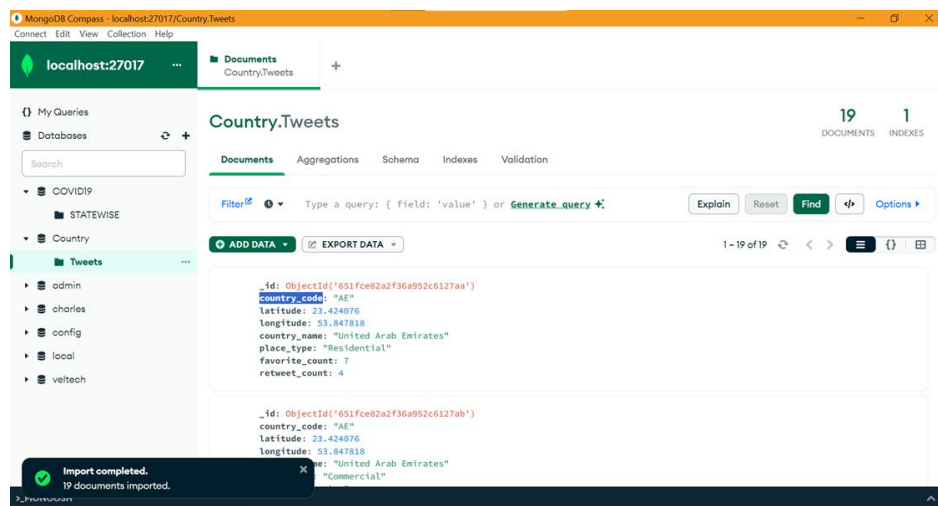
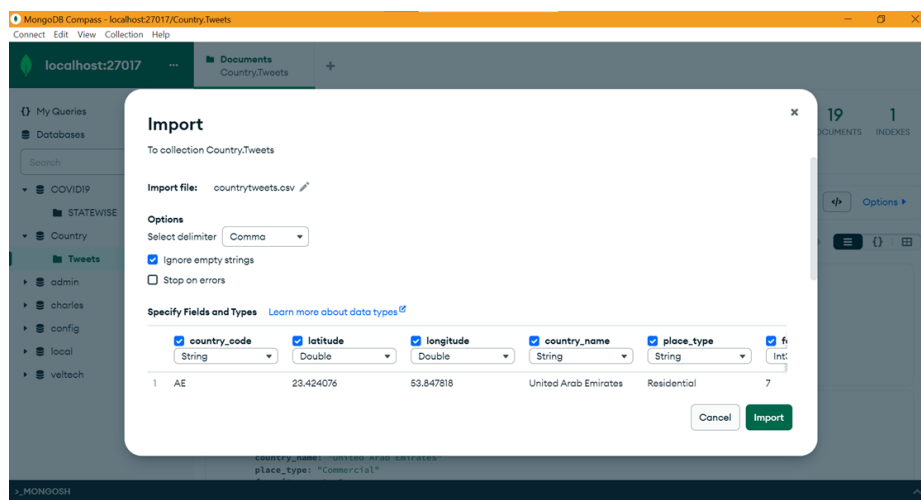
Total_retweet_count: 6

else:

```
print("No tweets found in the collection.")
```

```
# Close the MongoDB connection
```

```
client.close()
```




```
total_retweet_count = result[1]["total_retweet_count"]
print(f"Total favorite_count: {total_favorite_count}")
print(f"Total_retweet_count: {total_retweet_count}")
```

SAMPLE DATASET:

country_code	latitude	longitude	country_name	place_type	favorite_count	retweet_count
AE	23.424076	53.847818	United Arab Emirates	Residential	7	4
AE	23.424076	53.847818	United Arab Emirates	Commercial	5	2
AR	-38.416097	-63.616672	Argentina	Educational	4	6
AR	-38.416097	-63.616672	Argentina	Cultural	5	5
IN	20.593684	78.96288	India	Religious	80	3
IN	20.593684	78.96288	India	Historical	100	6
TH	15.870032	100.992541	Thailand	Shopping	2	5
IN	20.593684	78.96288	India	Sports	55	3
AO	-11.202692	17.873887	Angola	Residential	2	6
AQ	-75.250973	-0.071389	Antarctica	Commercial	4	4
AR	-38.416097	-63.616672	Argentina	Educational	5	9
AS	-14.270972	-170.132217	American Samoa	Cultural	3	6
BD	23.684994	90.356331	Bangladesh	Religious	0	4
AU	-25.274398	133.775136	Australia	Historical	8	4
AW	12.52111	-69.968338	Aruba	Shopping	2	8
AZ	40.143105	47.576927	Azerbaijan	Sports	4	5
EG	26.820553	30.802498	Egypt	Tourist	2	2
IN	20.593684	78.96288	India	Tourist	8	5
BD	23.684994	90.356331	Bangladesh	Religious	9	2

RESULT:



B. Find out top 10 most frequent topic words of the entire tweet message texts of your collection after lemmatization/stemming and removing all the Stop Words.

AIM:

PROCEDURE:

1. Install NLTK data in the PyCharm IDE
2. Import the NLTK libraries WordNetLemmatizer, PorterStemmer, stopwords
3. Read the sample tweet message from MongoDB or assign tweet = 'sample message'
4. Tokenize the tweet into words
5. Initialize lemmatizer and stemmer
6. Lemmatize and stem each word
7. Remove stop words
8. Define the list of English stop words
9. Join the filtered words back into a sentence
10. Print the results

PROGRAM:

```
import nltk
from nltk.stem import WordNetLemmatizer
from nltk.stem import PorterStemmer
from nltk.corpus import stopwords

nltk.download('punkt')
nltk.download('wordnet')
nltk.download('stopwords')

tweet = "The quick brown foxes are jumping over the lazy dogs' bones."

words = nltk.word_tokenize(tweet)
lemmatizer = WordNetLemmatizer()
stemmer = PorterStemmer()
lemmatized_words = [lemmatizer.lemmatize(word) for word in words]
stemmed_words = [stemmer.stem(word) for word in words]
words = nltk.word_tokenize(tweet)

stop_words = set(stopwords.words('english'))
filtered_words = [word for word in words if word.lower() not in stop_words]
filtered_text = ' '.join(filtered_words)
print("Original words:", words)
```


OUTPUT:

[nltk_data] Downloading package punkt to

[nltk_data] C:\Users\Lenovo\AppData\Roaming\nltk_data...

[nltk_data] Package punkt is already up-to-date!

[nltk_data] Downloading package wordnet to

[nltk_data] C:\Users\Lenovo\AppData\Roaming\nltk_data...

[nltk_data] Package wordnet is already up-to-date!

[nltk_data] Downloading package stopwords to

[nltk_data] C:\Users\Lenovo\AppData\Roaming\nltk_data...

[nltk_data] Package stopwords is already up-to-date!

Original words: ['The', 'quick', 'brown', 'foxes', 'are', 'jumping', 'over', 'the', 'lazy', 'dogs', '', 'bones', '.']

Lemmatized words: ['The', 'quick', 'brown', 'fox', 'are', 'jumping', 'over', 'the', 'lazy', 'dog', '', 'bone', '.']

Stemmed words: ['the', 'quick', 'brown', 'fox', 'are', 'jump', 'over', 'the', 'lazi', 'dog', '', 'bone', '.']

Tweet without stop words: quick brown foxes jumping lazy dogs bones .

```
print("Lemmatized words:", lemmatized_words)
print("Stemmed words:", stemmed_words)
print("Tweet without stop words:", filtered_text)
```

RESULT: